

AEROSPIKE

Natalie Pistunovich

@NataliePis @aerospikeDB





The **aerospike engine** is a type of rocket engine that maintains its aerodynamic efficiency across a wide range of altitudes.

It belongs to the class of altitude compensating nozzle engines. A vehicle with an aerospike engine uses 25-30% less fuel at low altitudes, where most missions have the greatest need for thrust.

Aerospike engines have been studied for a number of years and are the baseline engines for many single-stage-to-orbit designs and Space Shuttle main engine.

- wikipedia







Terminology

SQL vs. NoSQL

Parameter	SQL	NOSQL
Definition	SQL databases are primarily called RDBMS or Relational Databases	NoSQL databases are primarily called as Non-relational or distributed database
Design for	Traditional RDBMS uses SQL syntax and queries to analyze and get the data for further insights. They are used for OLAP systems.	NoSQL database system consists of various kind of database technologies. These databases were developed in response to the demands presented for the development of the modern application.
Query Language	Structured query language (SQL)	No declarative query language
Type	SQL databases are table based databases	NoSQL databases can be document based, key-value pairs, graph databases

SQL vs. NoSQL

Parameter	SQL	NOSQL
Schema	SQL databases have a predefined schema	NoSQL databases use dynamic schema for unstructured data.
Ability to scale	SQL databases are vertically scalable	NoSQL databases are horizontally scalable
Hierarchical data storage	SQL databases are not suitable for hierarchical data storage.	More suitable for the hierarchical data store as it supports key-value pair method.
Variations	One type with minor variations.	Many different types which include key-value stores, document databases, and graph databases.
Development Year	It was developed in the 1970s to deal with issues with flat file storage	Developed in the late 2000s to overcome issues and limitations of SQL databases.

When to NoSQL?

- Low-latency, low-overhead API to access data
- A few database features with high scale
- Avoid data/schema pre-design altogether for simple applications and agility
- Fast lookup by primary key

CAP

CAP Theorem is a concept that a distributed database system can only have 2 of the 3:

1. Consistency
2. Availability
3. Partition Tolerance

1. Consistency

All nodes see the same data at the same time.

Performing a read operation will return the value of the most recent write operation causing all nodes to return the same data.

A system has consistency if a transaction starts with the system in a consistent state, and ends with the system in a consistent state.

In this model, a system can (and does) shift into an inconsistent state during a transaction, but the entire transaction gets rolled back if there is an error during any stage in the process.

2. Availability

Every request gets a response on success/failure.

Achieving availability in a distributed system requires that the system remains operational 100% of the time.

Every client gets a response, regardless of the state of any individual node in the system.

This metric is trivial to measure: either you can submit read/write commands, or you cannot. Hence, the databases are time independent as the nodes need to be available online at all times.

3. Partition Tolerance

The system continues to run, despite the number of messages being delayed by the network between nodes.

A system that is partition-tolerant can sustain any amount of network failure that doesn't result in a failure of the entire network.

Data records are sufficiently replicated across combinations of nodes and networks to keep the system up through intermittent outages.

ACID

The principles that database transactions should adhere to, to ensure that data doesn't become corrupt as a result of a failure

1. Atomicity
2. Consistency
3. Isolation
4. Durability

Atomicity

Means that you guarantee that either all of the transaction succeeds or none of it does. You don't get part of it succeeding and part of it not. If one part of the transaction fails, the whole transaction fails. With atomicity, it's all or nothing.

Consistency

Guarantees that all data will be consistent. All data will be valid according to all defined rules, including any constraints, cascades, and triggers that have been applied on the database.

Isolation

Guarantees that all transactions will occur in isolation. No transaction will be affected by any other transaction. So a transaction cannot read data from any other transaction that has not yet completed.

Durability

Means that once a transaction is committed, it will remain in the system - even if there's a system crash immediately following the transaction. Any changes from the transaction must be stored permanently. If the system tells the user that the transaction has succeeded, the transaction must have, in fact, succeeded.

Analyses

Over the past six years, Jepsen has analyzed over two dozen databases, coordination services, and queues—and we've found replica divergence, data loss, stale reads, read skew, lock conflicts, and much more. Here's every analysis we've published.

Aerospike	2015-05-04	3.5.4
	2018-03-07	3.99.0.3
Cassandra	2013-09-24	2.0.0
Chronos	2015-08-10	2.4.0
CockroachDB	2017-02-16	beta-20160829
Crate	2016-06-28	0.54.9
Dgraph	2018-08-23	1.0.2
Elasticsearch	2014-06-15	1.1.0
	2015-04-27	1.5.0
etcd	2014-06-09	0.4.1
	2020-01-30	3.4.3
FaunaDB	2019-03-05	2.5.4
Hazelcast	2017-10-06	3.8.3
Kafka	2013-09-24	0.8 beta
MariaDB Galera	2015-09-01	10.0

<https://jepsen.io>

Analyses

Over the past six years, Jepsen has analyzed over two dozen databases, coordination services, and queues—and we've found replica divergence, data loss, stale reads, read skew, lock conflicts, and much more. Here's every analysis we've published.

Aerospike	2015-05-04	3.5.4
	2018-03-07	3.99.0.3

“Aerospike does appear to provide linearizability through network partitions and process crashes”

- Kyle Kingsbury, Jepsen.io (<https://jepsen.io/analyses/aerospike-3-99-0-3>)

etcd	2015-04-27	1.5.0
	2014-06-09	0.4.1
	2020-01-30	3.4.3
FaunaDB	2019-03-05	2.5.4
Hazelcast	2017-10-06	3.8.3
Kafka	2013-09-24	0.8 beta
MariaDB Galera	2015-09-01	10.0

<https://jepsen.io>





About Aerospike

The Hummingbird

- Takes ~20ms to flap its wings, 50-70 wing flaps every second
- The smallest bird in the world, weigh less than a penny (2 grams)

AEROSPIKE

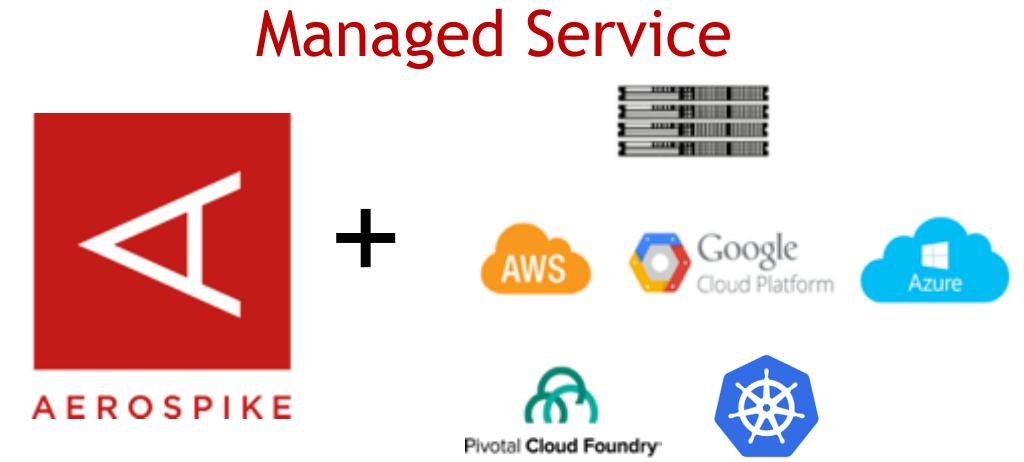
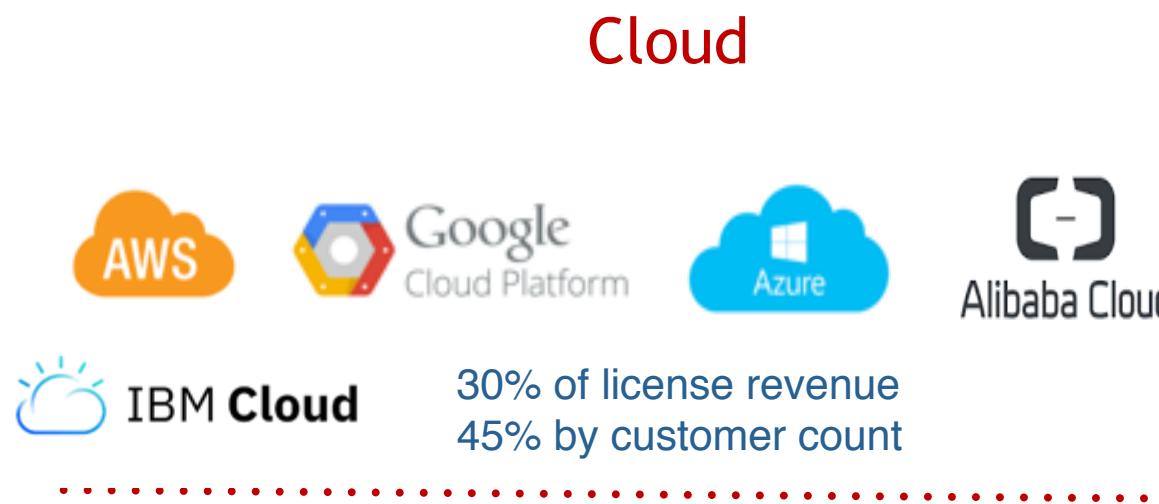


Aerospike

- ~300K read/writes every second, 99.9% are <1ms latency
- Unmatched reliability and uptime and deployable anywhere



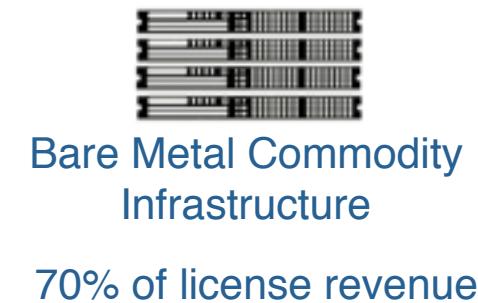
Deployment



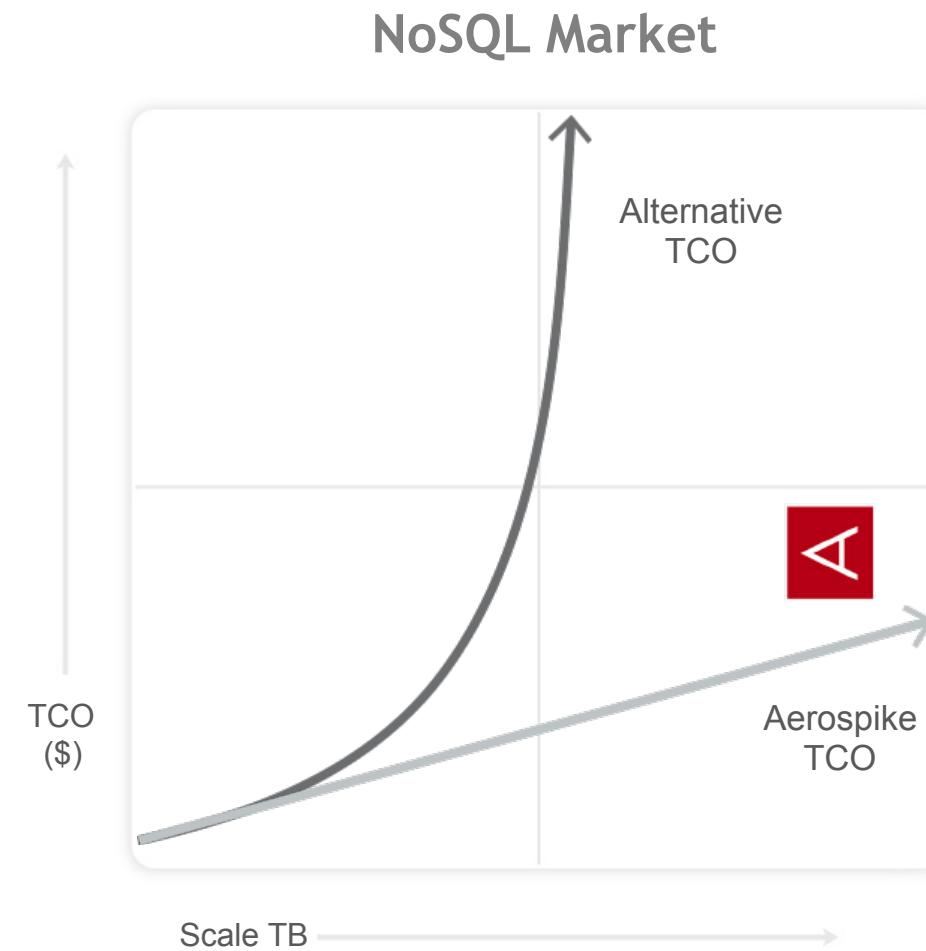
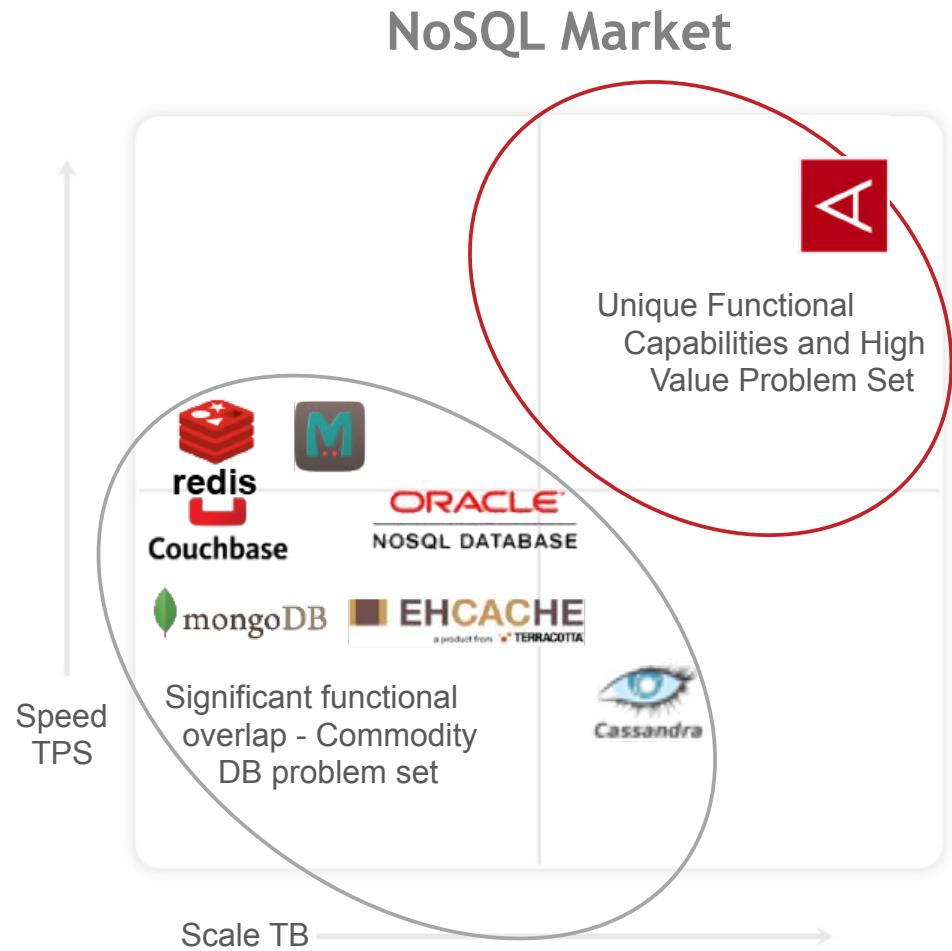
Orchestration Platforms



Traditional Linux Packages



Multiple NoSQL Technologies



Customers

ADTECH	ECOMMERCE	FINANCIAL	NEWTECH	TELCO	GAMING
 AppNexus  nielsen  Oath   AppsFlyer      	       	          	       	      	     



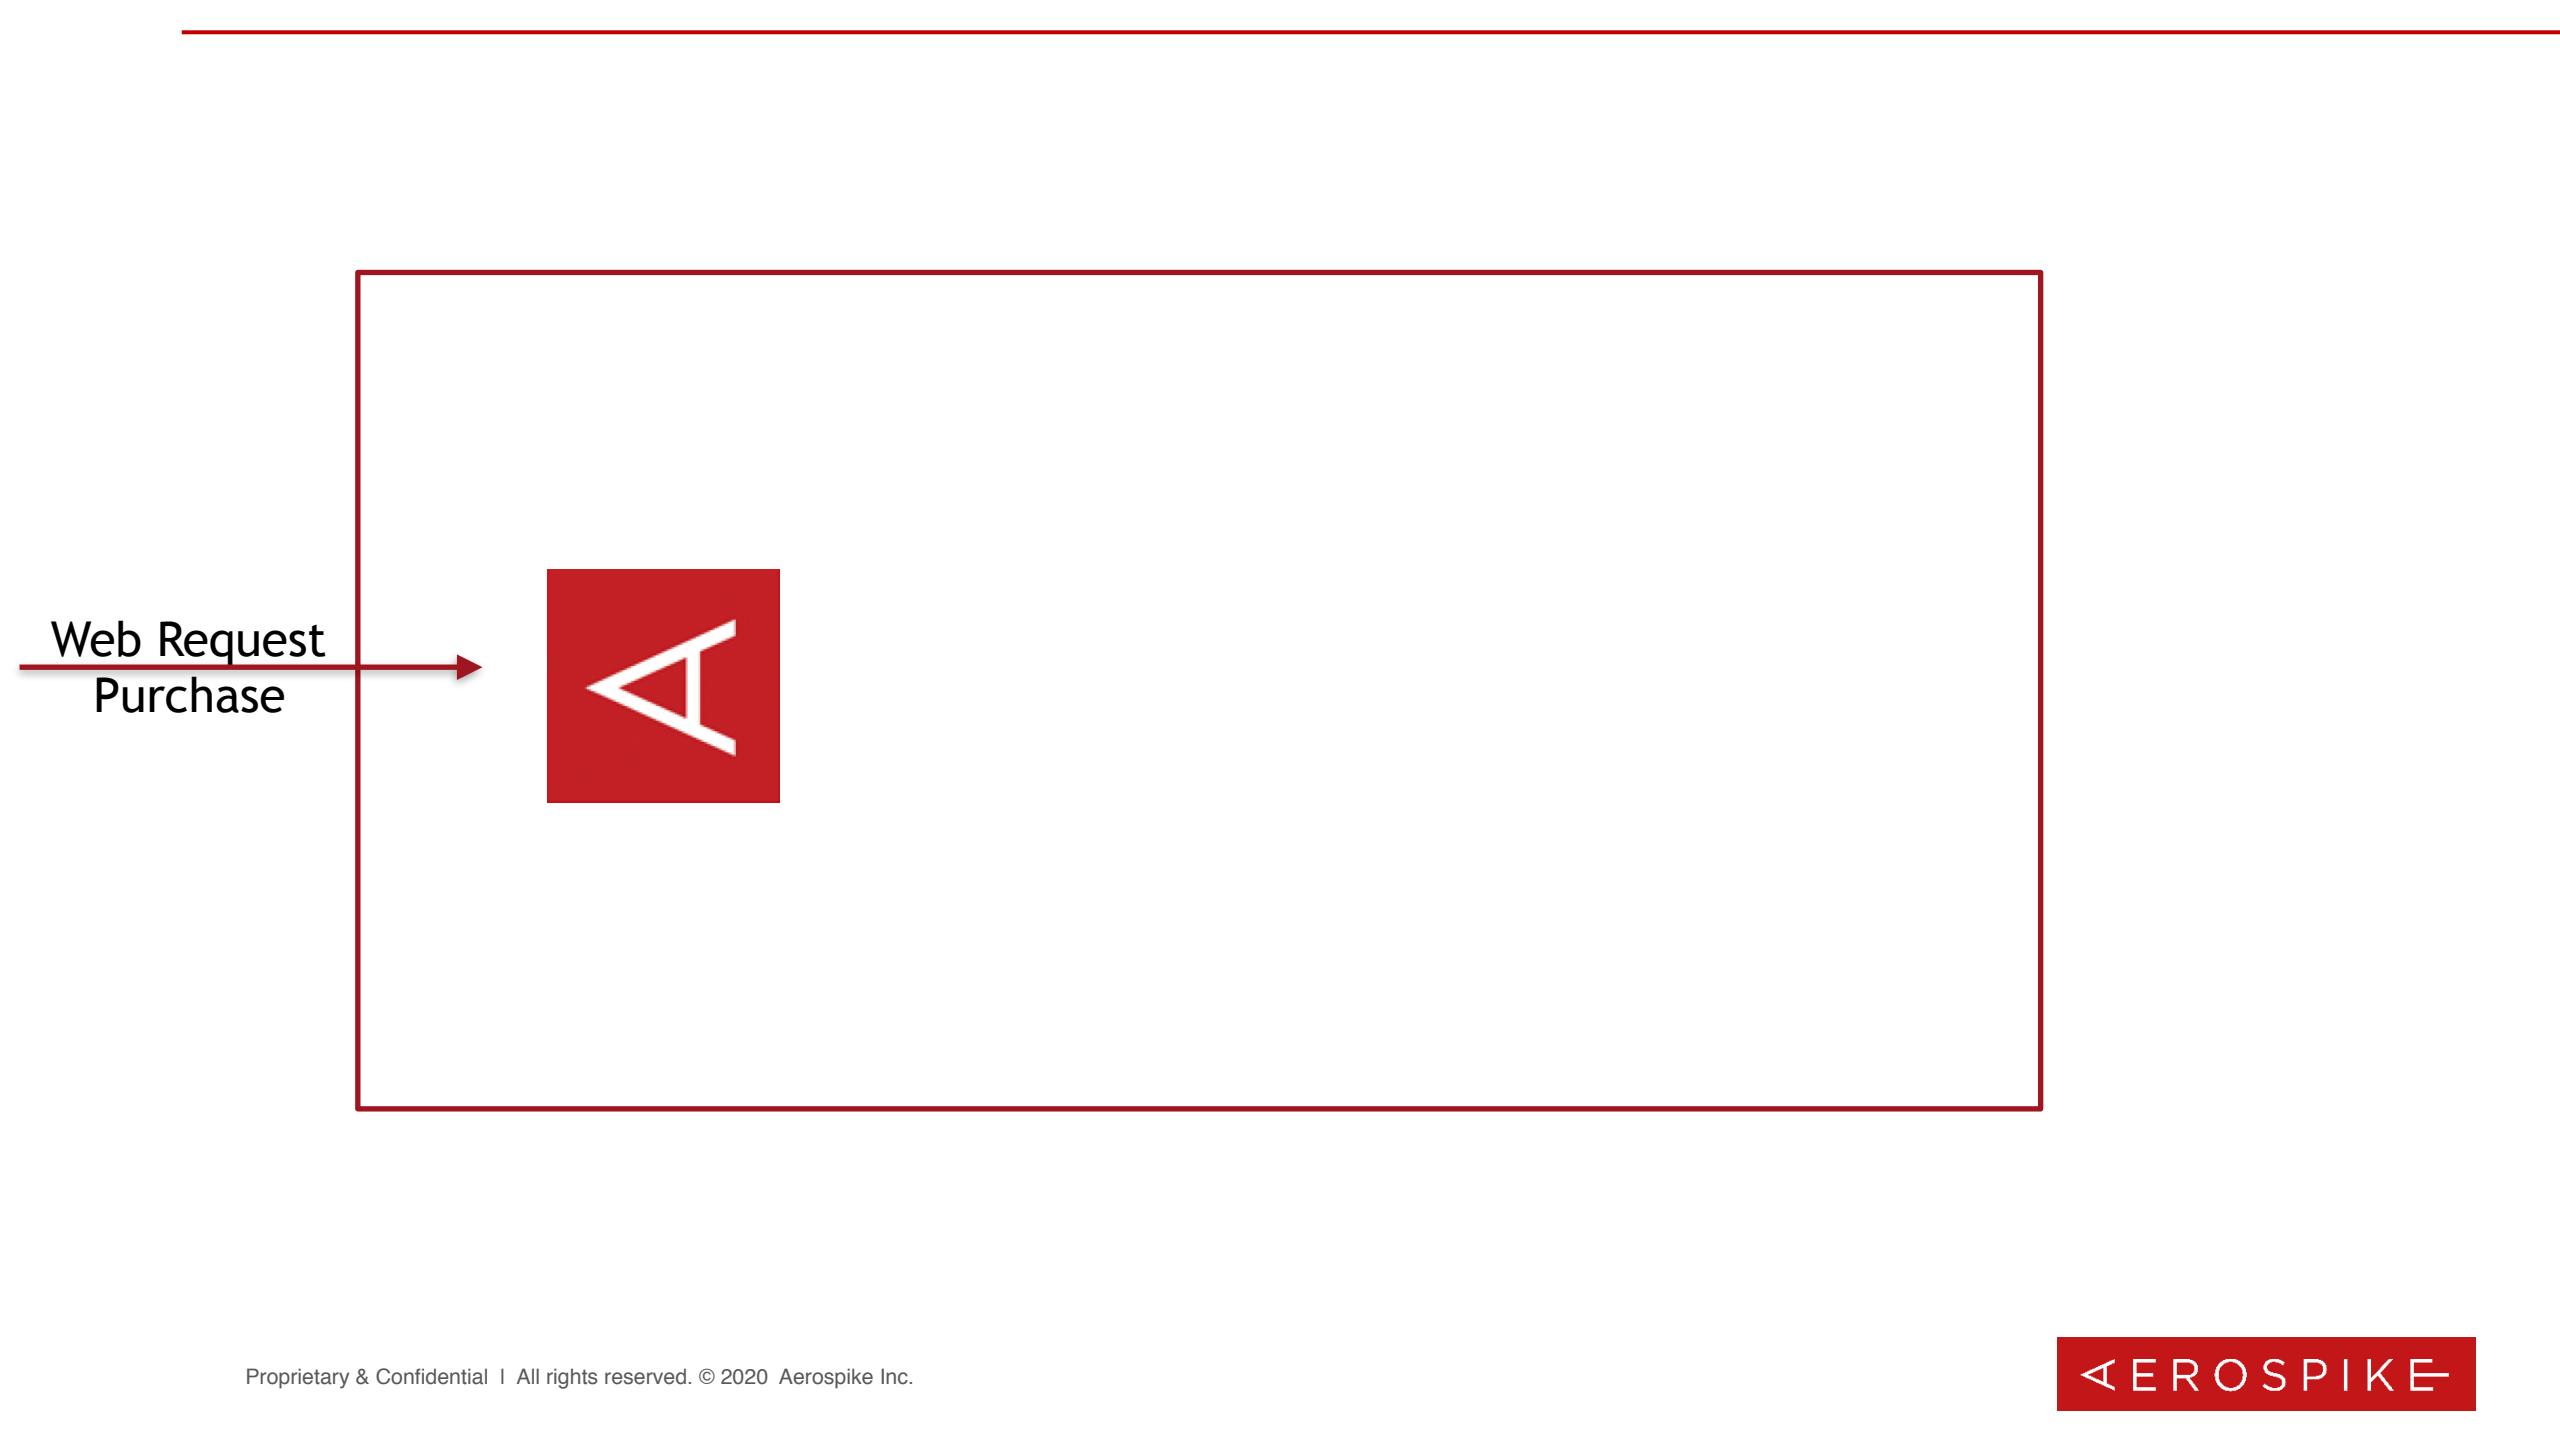
A Use Case

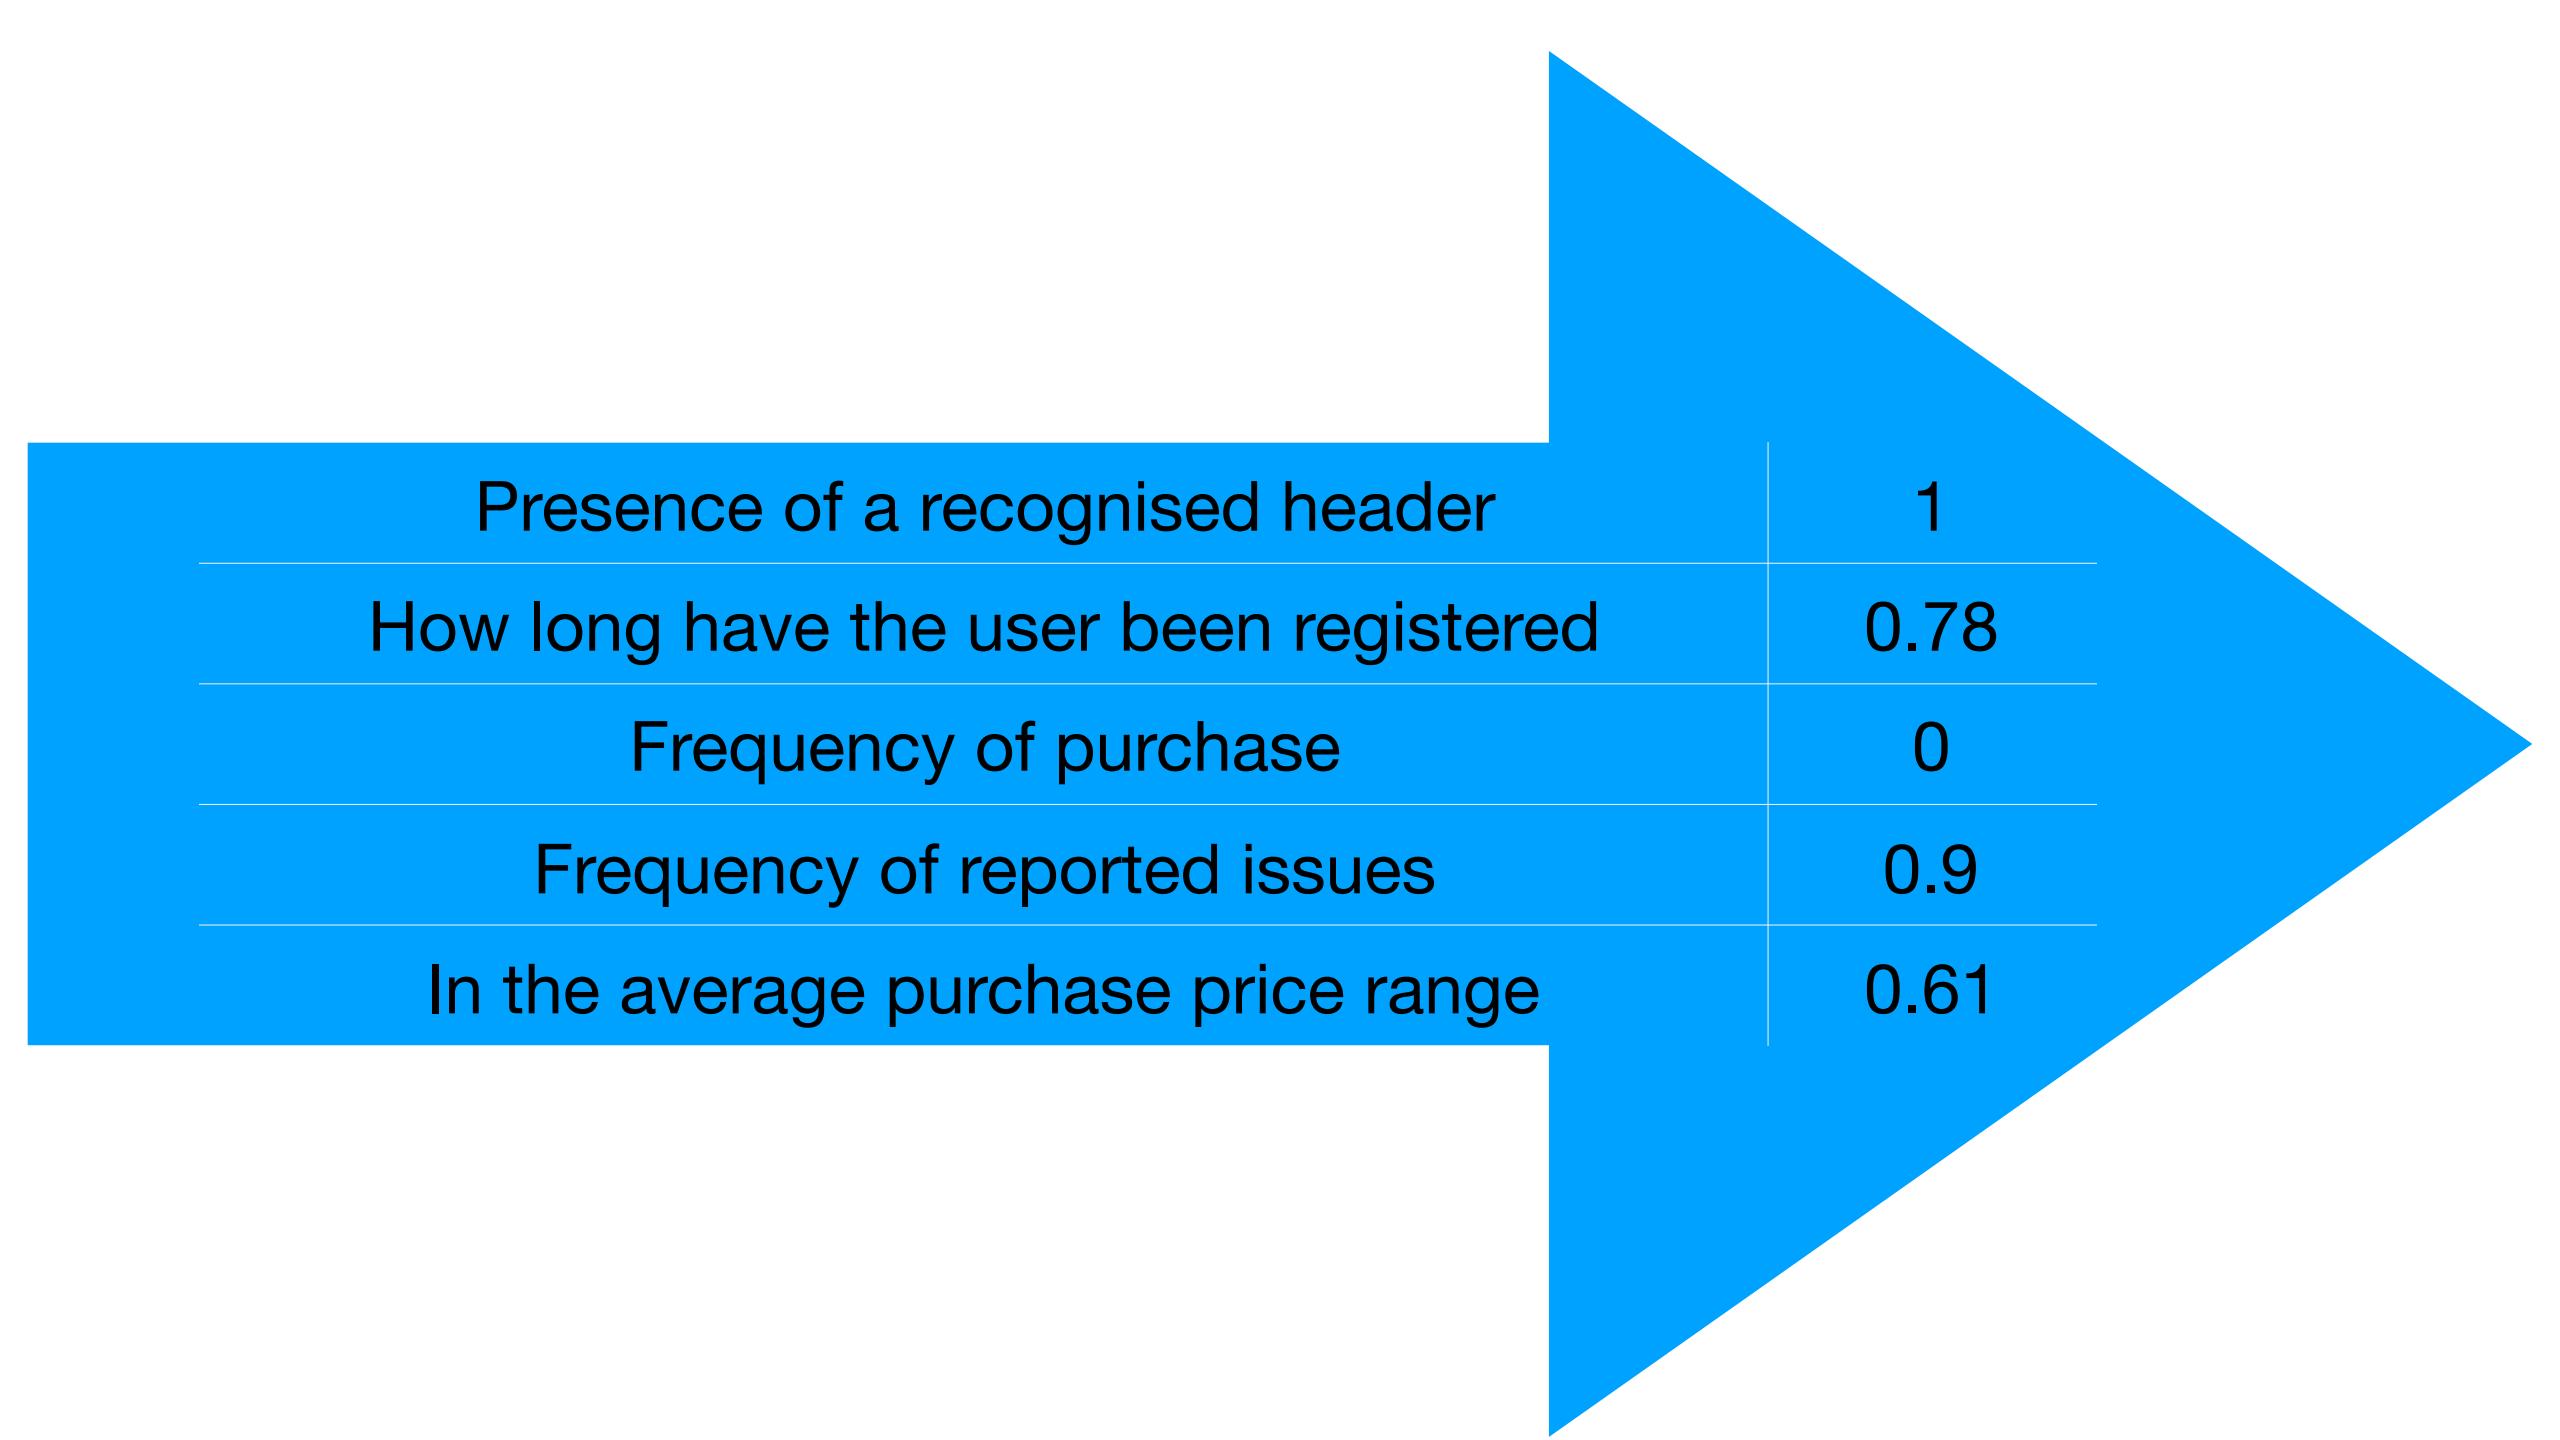
Business Challenge

- Tracking 150 million customers enrollment, payment and invoicing preferences, and profile data (e.g. IP address, location)
- Frequent user/device movement opens door for fraud
- Every payment transaction requires hundreds of DB reads/writes
- Missed latency SLA leads to lost business – fraud algorithm has a process window during payments
- Caching solution too expensive

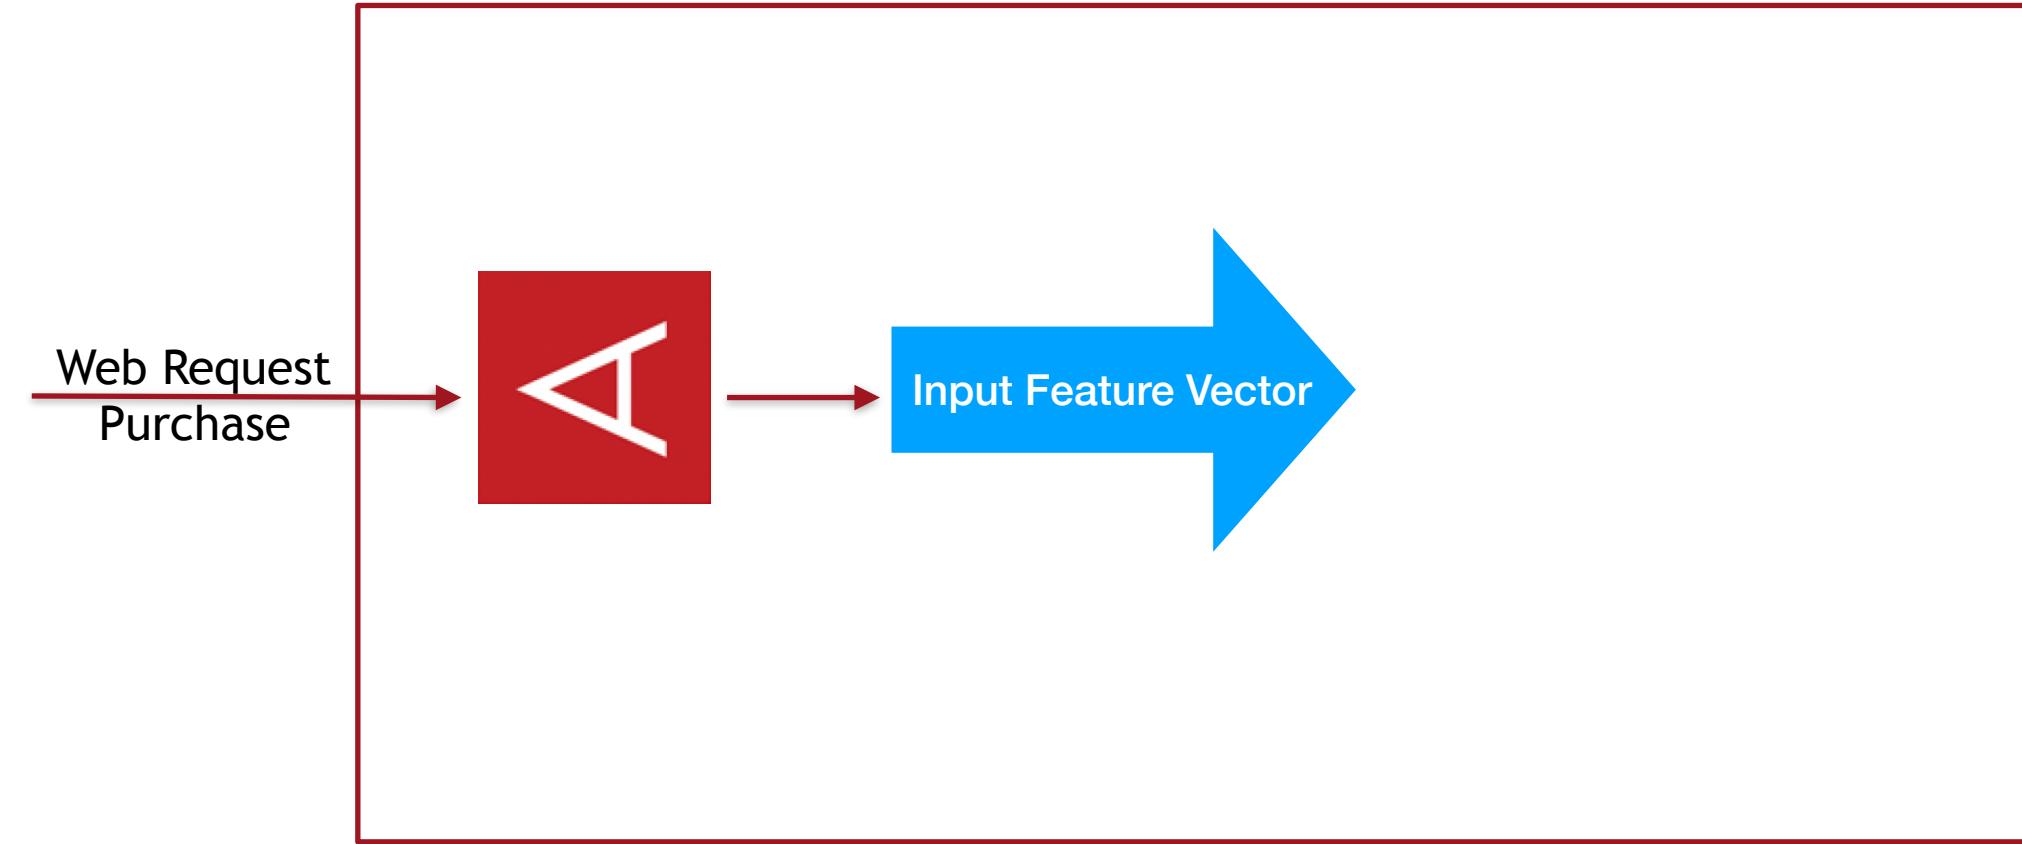
Need to scale up

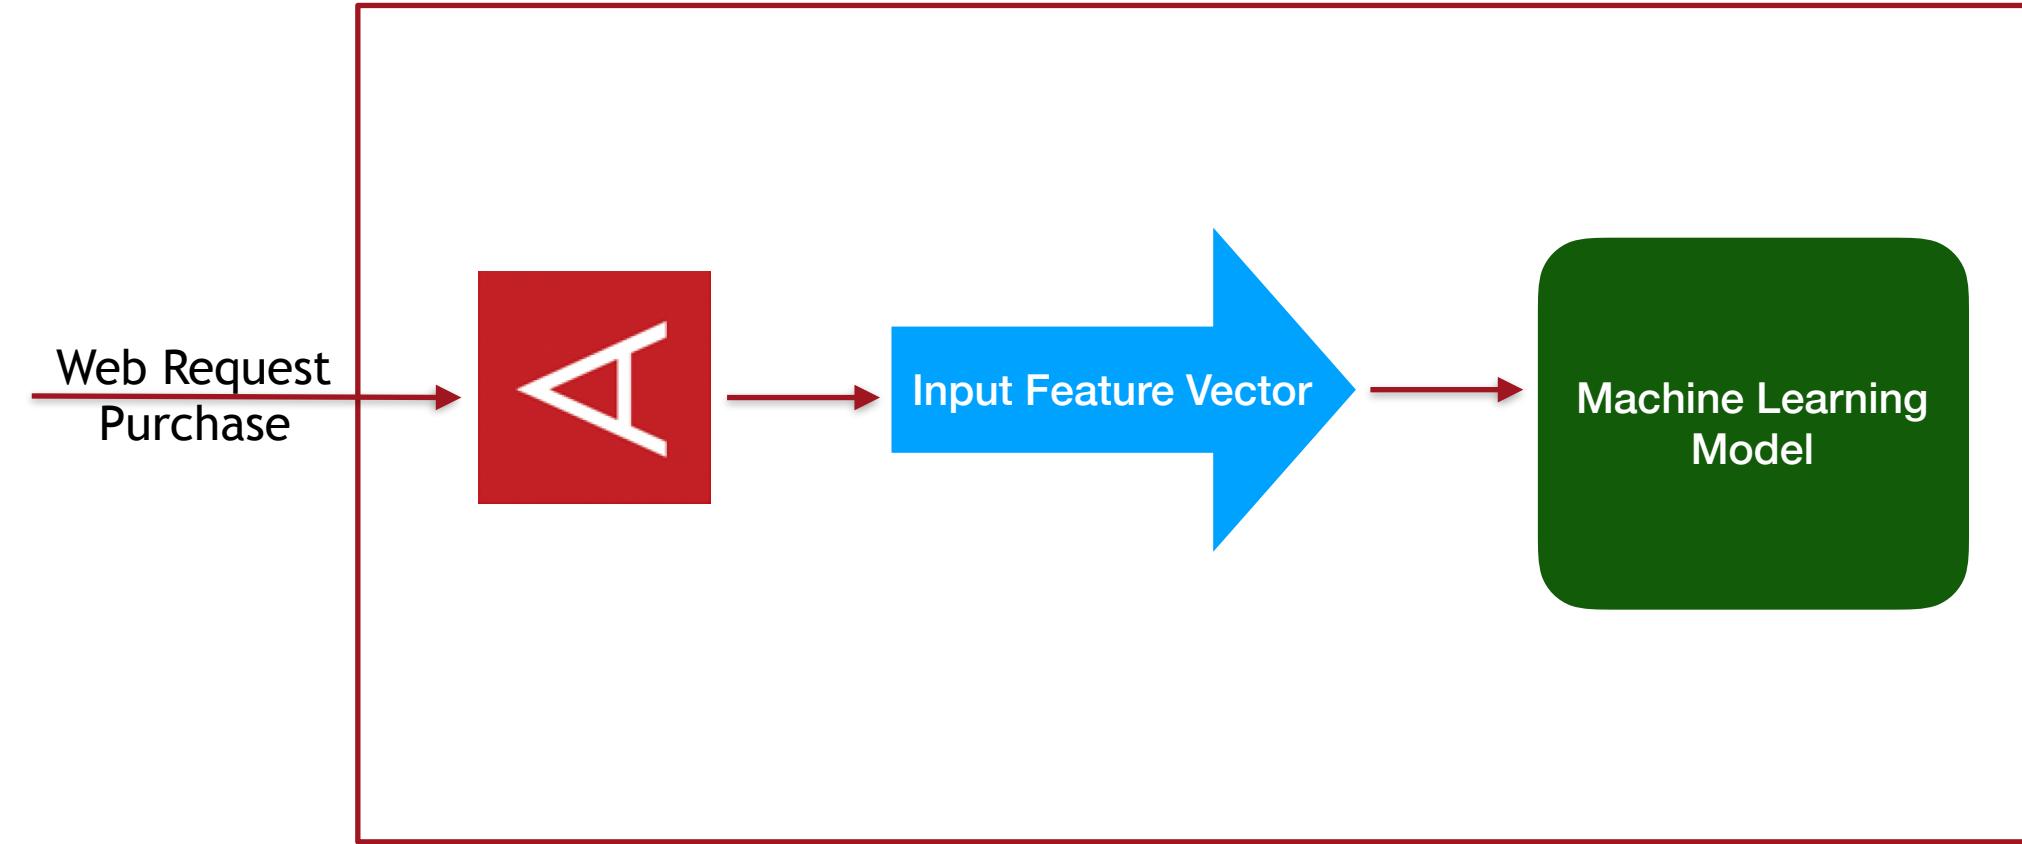
- 10 → 100 TB
- 10B → 100 B objects
- 200k → 1 Million+ TPS

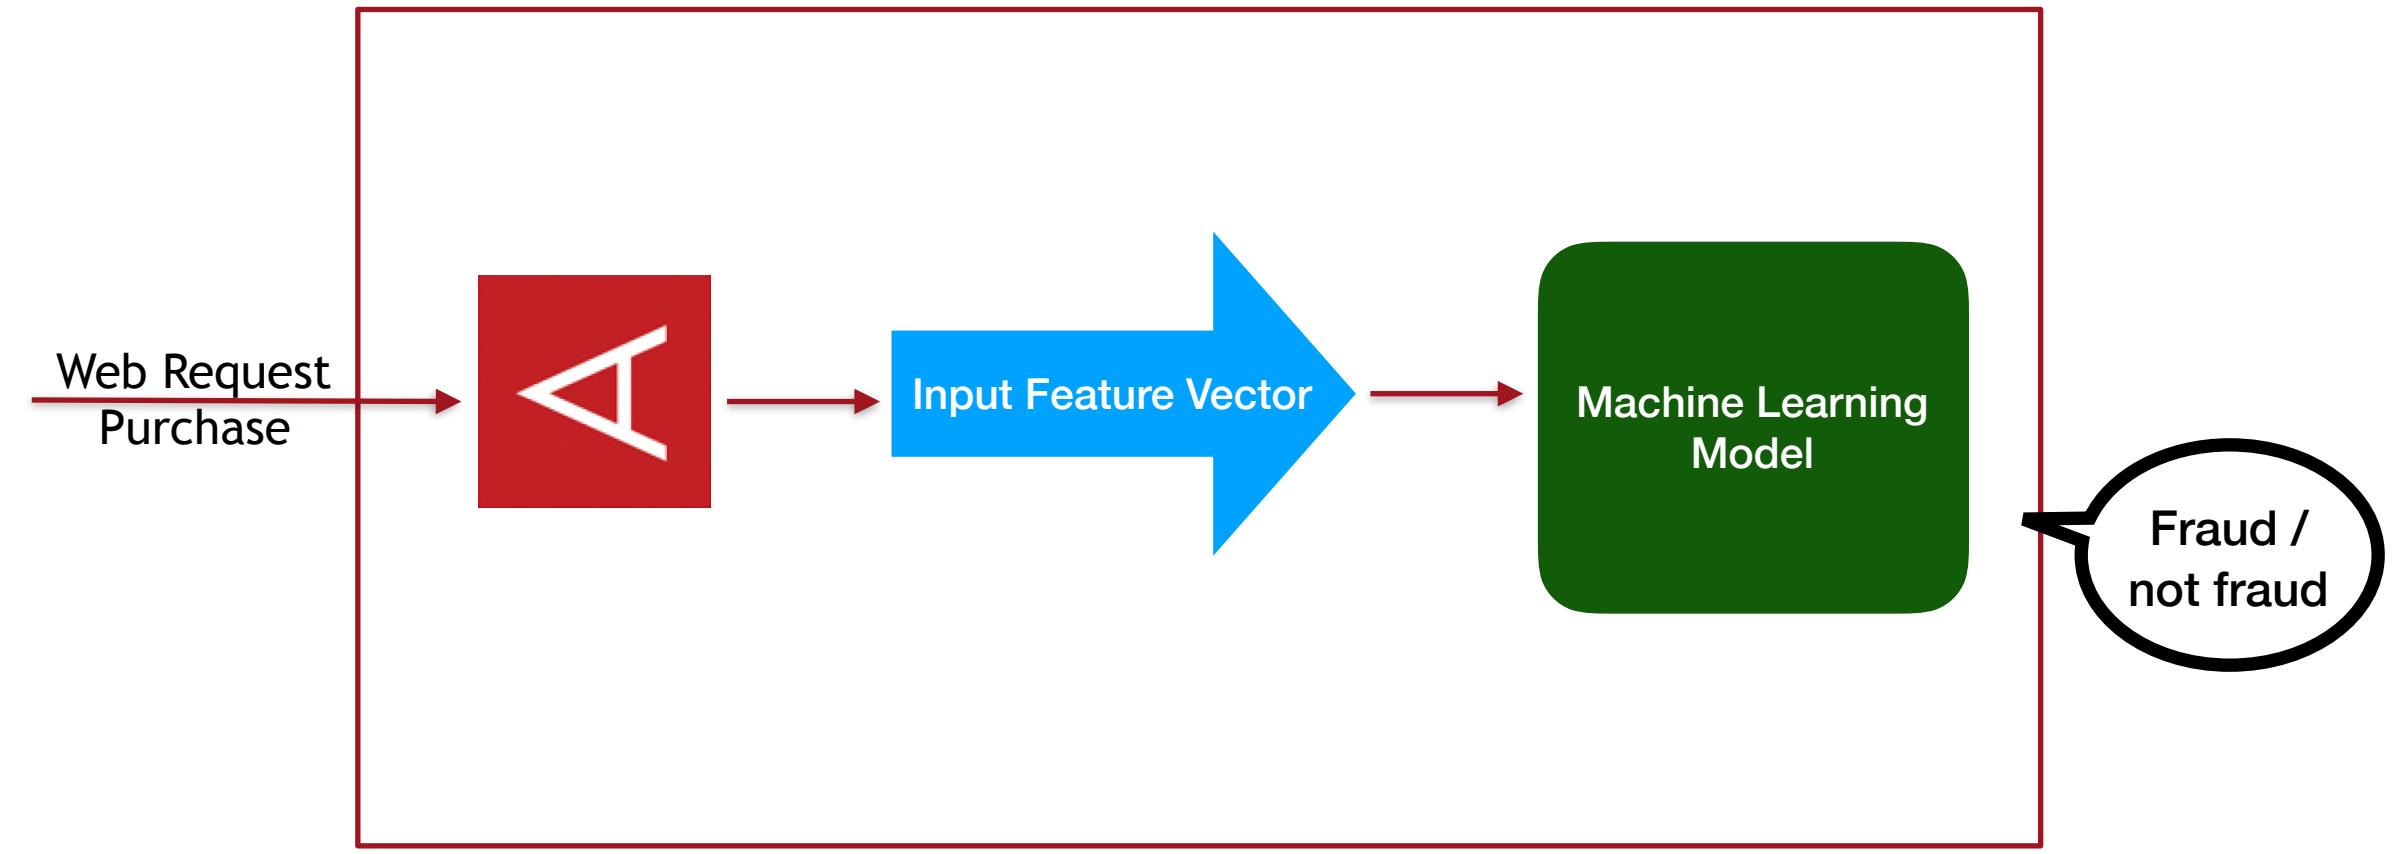


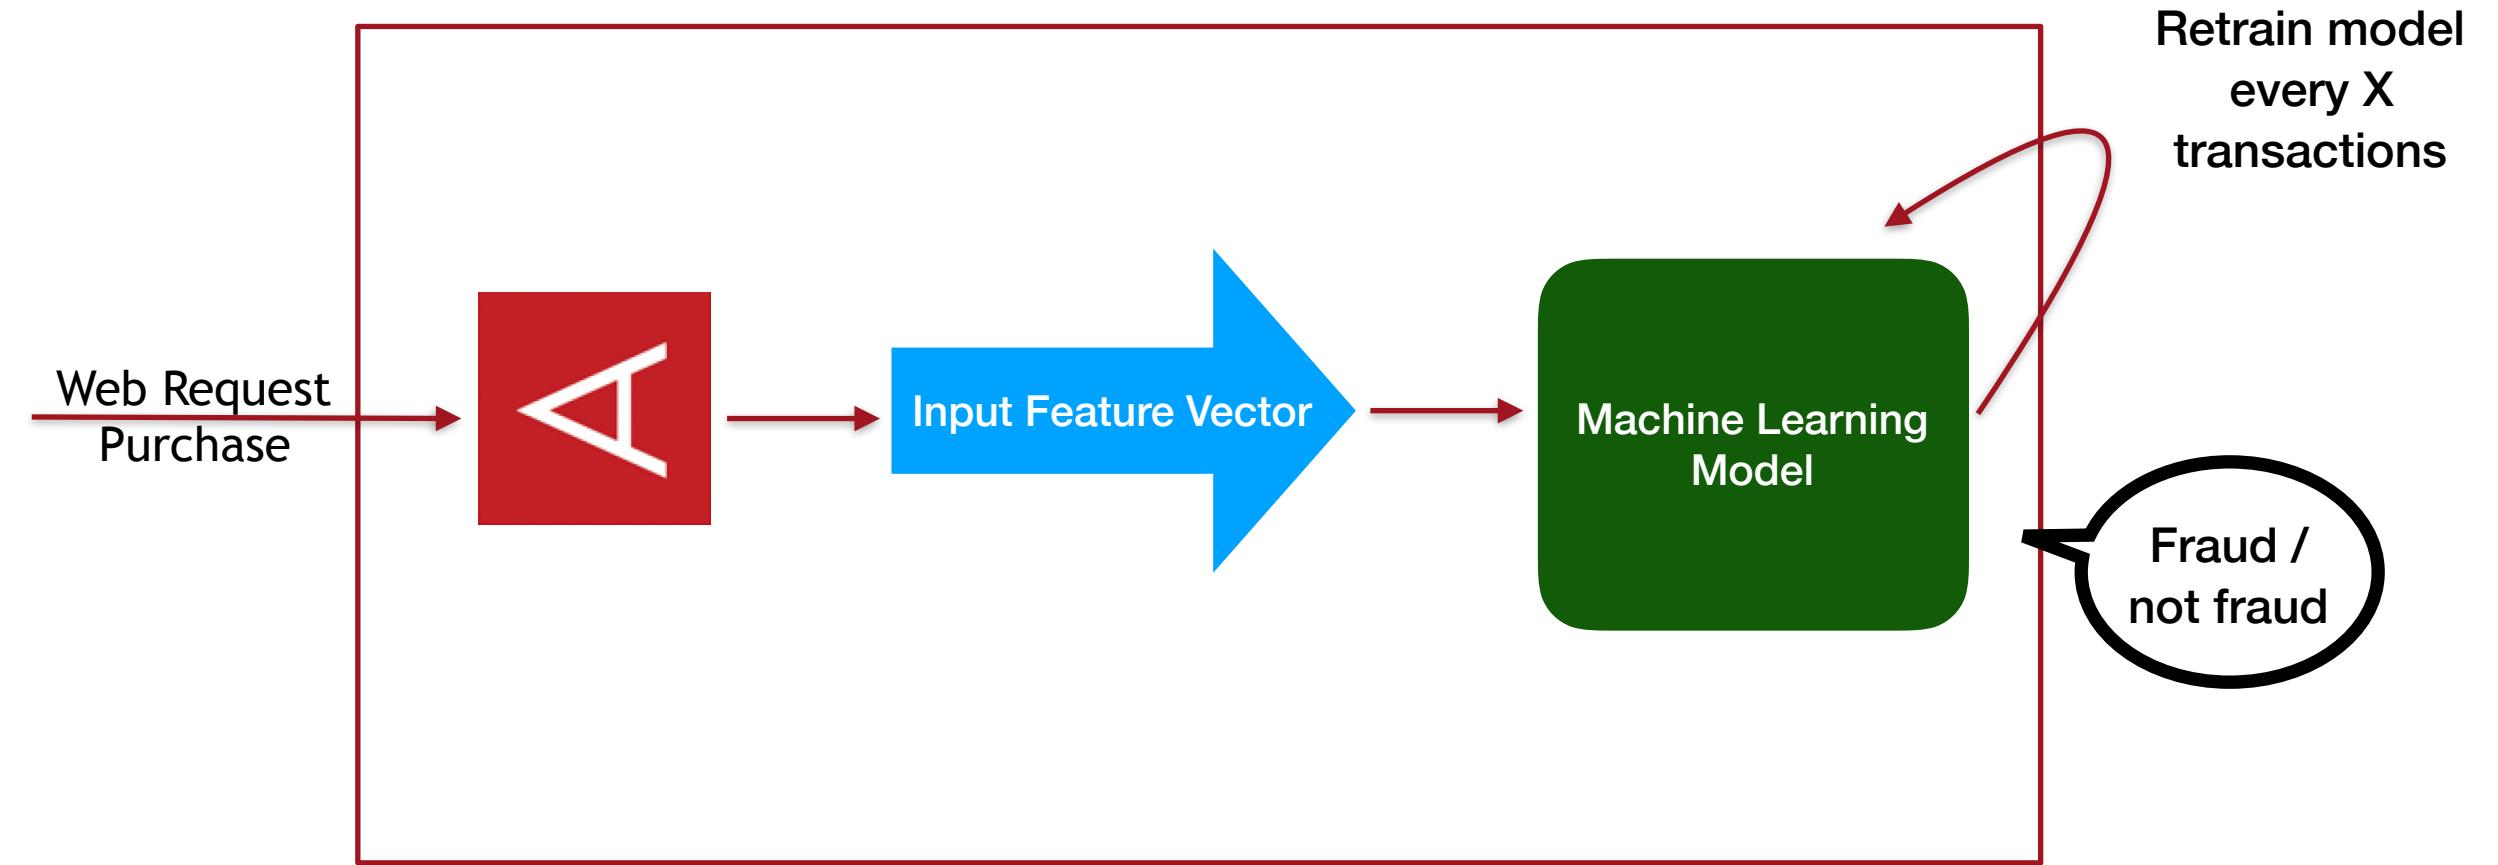


Presence of a recognised header	1
How long have the user been registered	0.78
Frequency of purchase	0
Frequency of reported issues	0.9
In the average purchase price range	0.61









Powering Global Fraud Prevention Network

✓ \$280B+ payments annually

Replaced Oracle + Terracota

✓ Reduced server footprint 15x

Improved SLAs

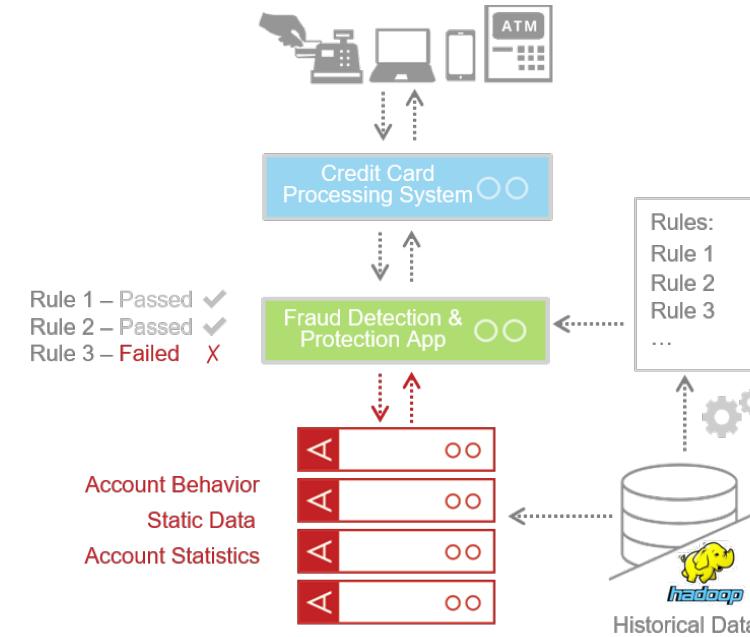
✓ 30x reduction in false positives

Increased revenue

✓ 10x improvement in fraud calculation data used

“PayPal is innovating deep analytics to rapidly respond to emerging fraud patterns, then deploying into an event-driven, fast data, in-memory architecture to accelerate detection, reduce losses and achieve near-continuous availability.”

Mikhail Kourjanski, PhD. - Lead Data Architect, Risk and Compliance Management Platform, PayPal





What's So Special About the Architecture?

Hybrid Memory Architecture



- Flash Optimized Storage Layer
- Multi-threaded & Massively Parallel
- Self-healing Clusters

Hybrid Memory Architecture



- **Flash Optimized Storage Layer**
- Multi-threaded & Massively Parallel
- Self-healing Clusters

Flash Optimized Storage Layer: Indexes in DRAM, Data on SSD

- Small amount of DRAM
- No cache so no cache misses
- Optimized for SSDs
- Continual Defragmentation
- Direct Device Access

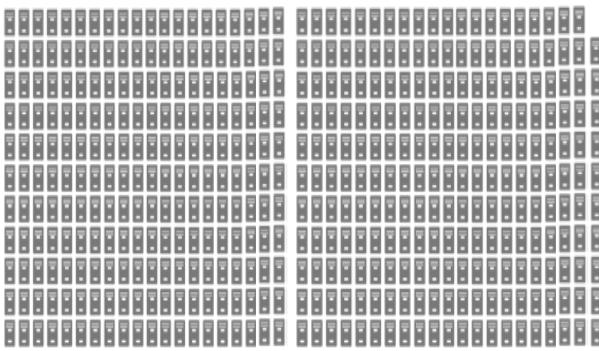


Total Cost of Ownership

AdTech Use Case - Signal

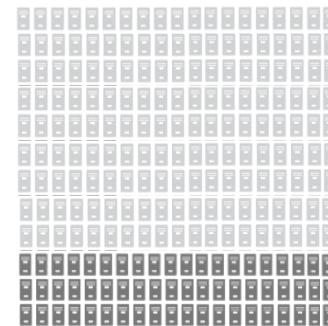
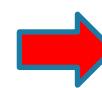
TCO savings of 1 app \$7.75M over 3 yrs

450 Cassandra Servers



cassandra

Only 60 Aerospike Servers



AEROSPIKE

	Year 1 (\$M USD)	Year 2 (\$M USD)	Year 3 (\$M USD)	Total (\$M USD)
Cassandra	\$2.39	\$3.29	\$4.53	\$10.21
Aerospike	\$1.69	\$.24	\$.52	\$2.46
Total OpEx Savings from Aerospike (\$M USD)	\$.70	\$3.05	\$4.01	\$7.75

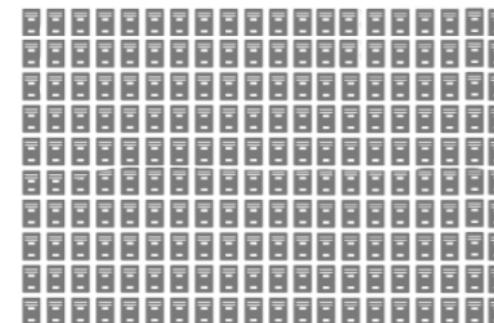
"Free is even too expensive compared to Aerospike."
- Jason Yanowitz, EVP and CTO, Signal

Gaming Use Case - Playtika



TCO savings of app \$1.4M/year & \$4.2M over 3 yrs

200 Couchbase Servers



Couchbase

DRAM & SSD

200K/s Reads+Writes

Only 30 Aerospike Servers



AEROSPIKE

SSD & DRAM

600K/s Reads+Writes

"Simplicity of IT Ops frees teams up for future-looking initiatives."
- Guy Almog, Director of Engineering, Playtika

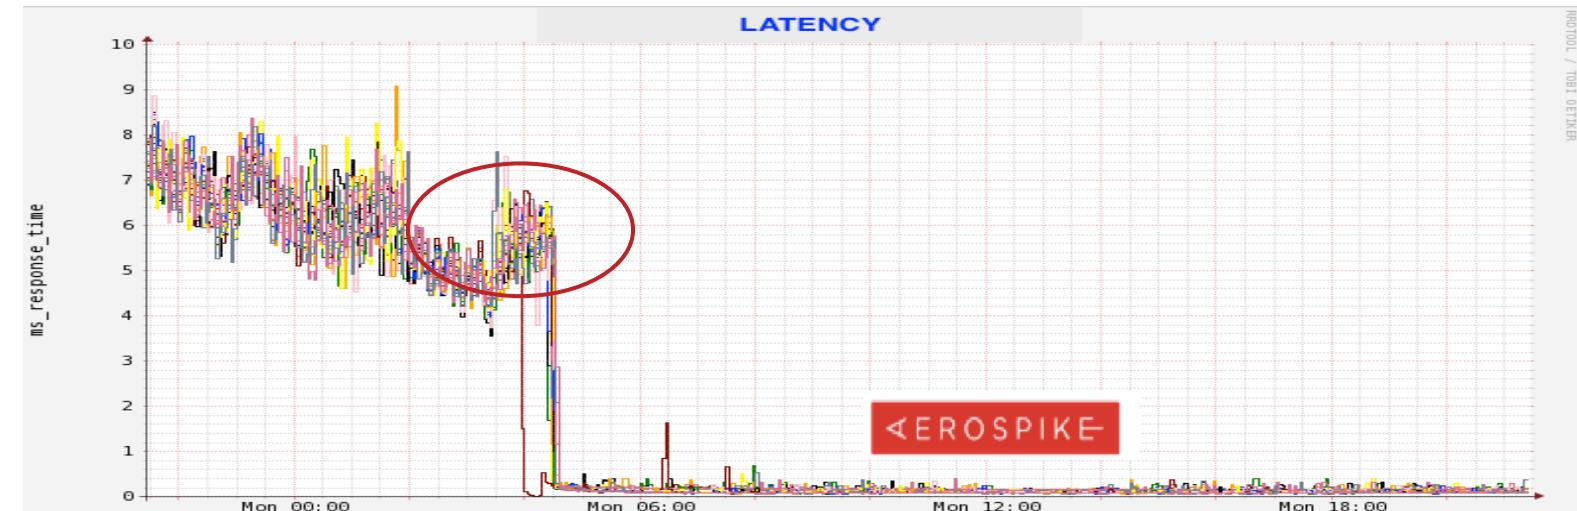
Aerospike's cacheless architecture is simpler, faster and cheaper than cache-based architectures

AEROSPIKE

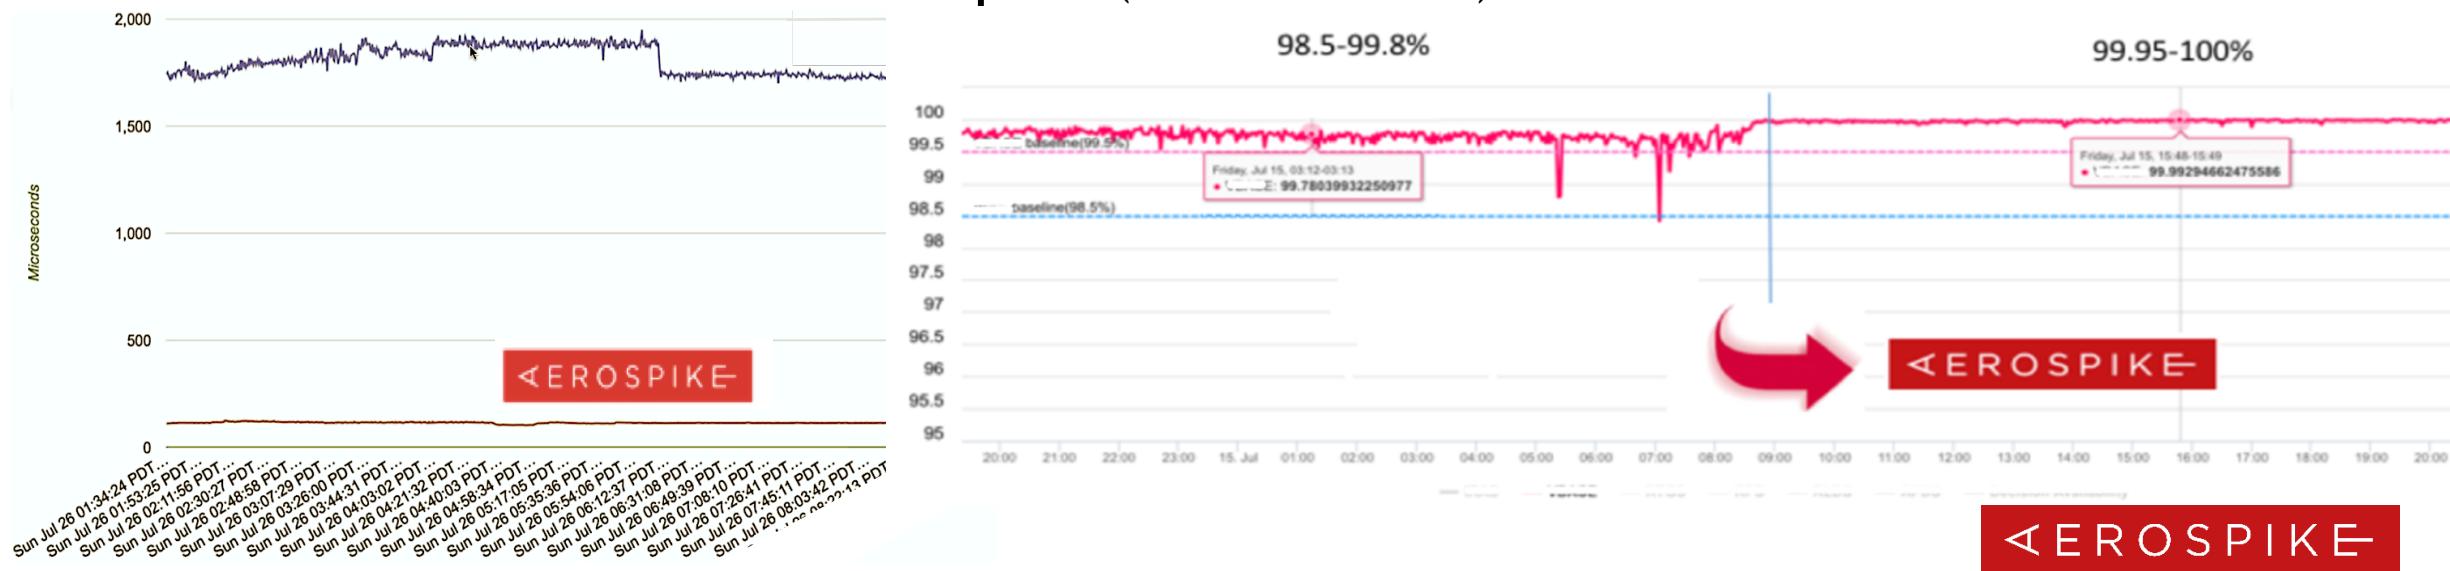
Actual Customer Comparisons when moving to Aerospike

Cassandra to Aerospike

- Digital identity management



Relational + Cache to Aerospike (Fraud detection)



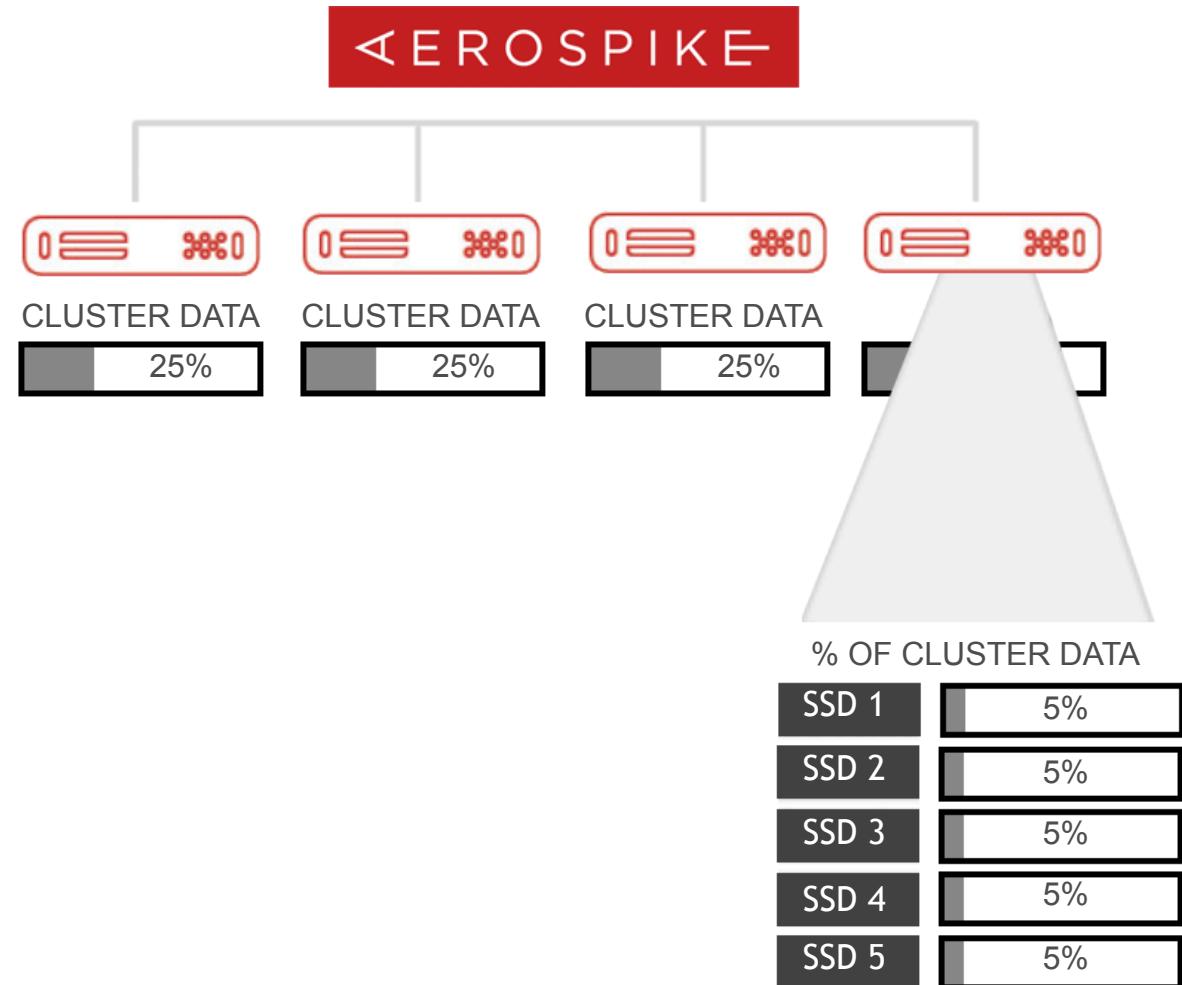
Hybrid Memory Architecture



- Flash Optimized Storage Layer: Indexes in DRAM, Data on SSD
- **Multi-threaded & Massively Parallel**
- Self-healing Clusters

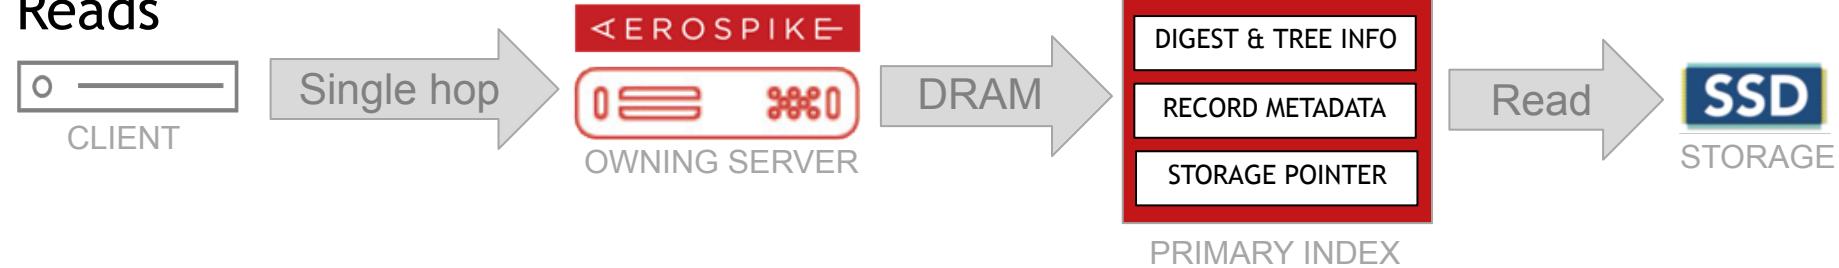
Multi-threaded & Massively Parallel

- Linear scaling
- Automatic Distribution of Data Using
- Smart Clients

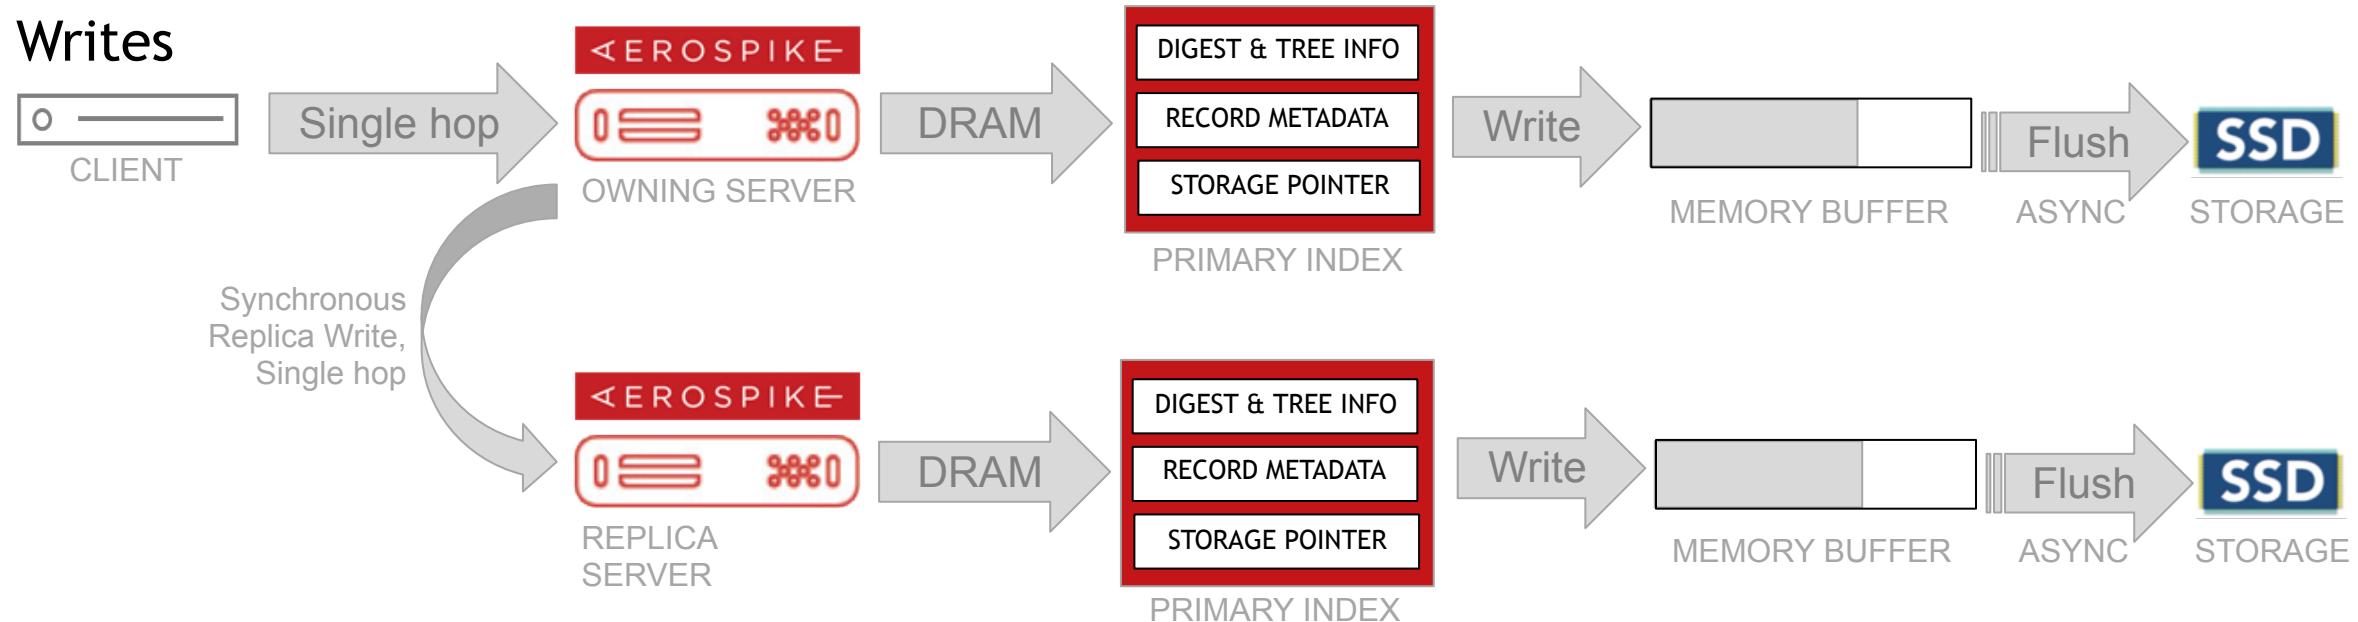


Smart Clients

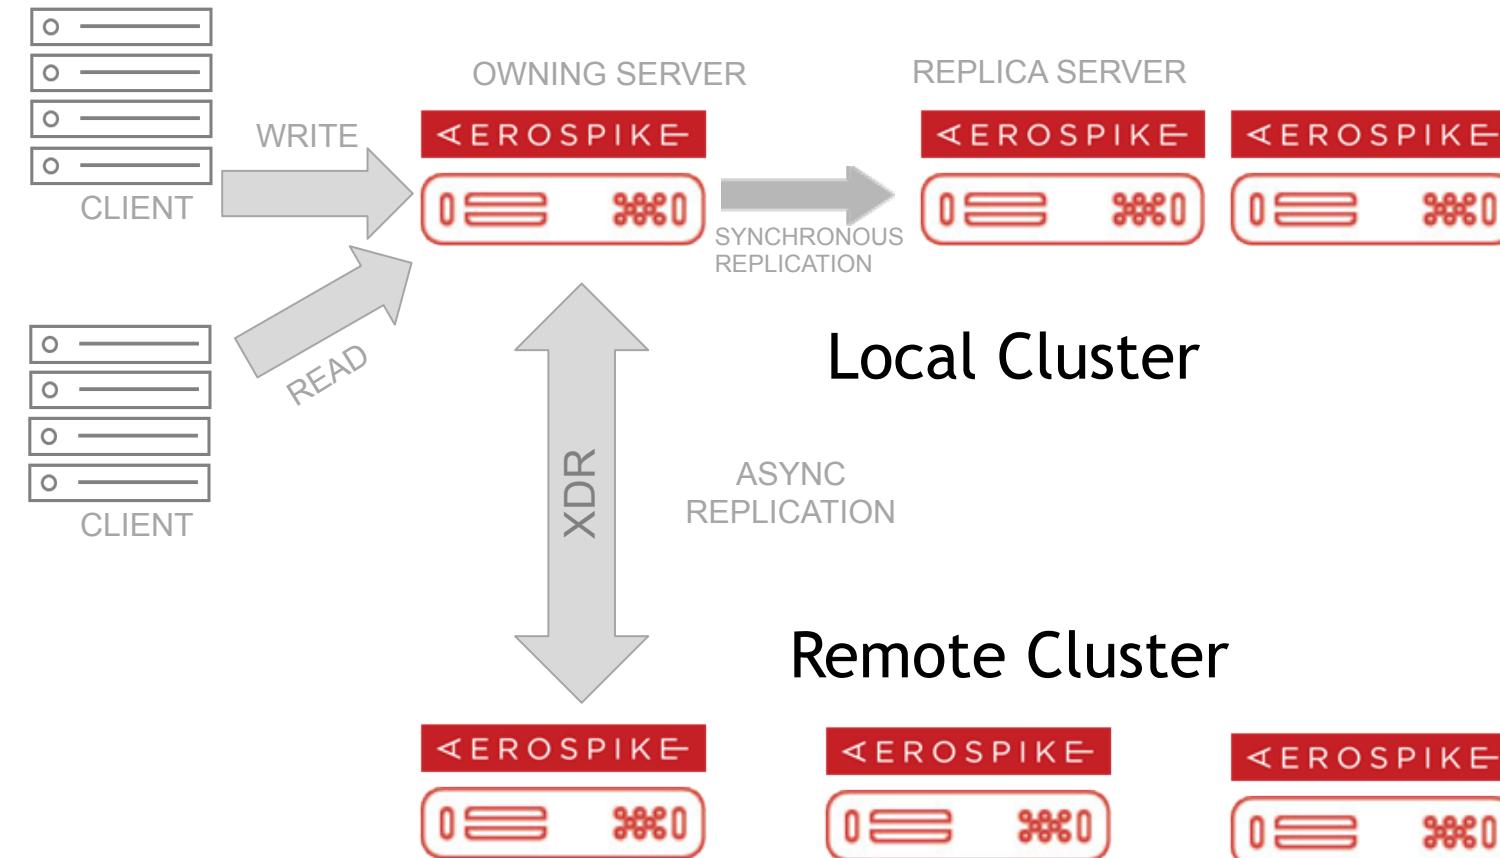
Reads



Writes



Data Correctness



Cacheless Architecture

- ✓ Reduce server count
- ✓ No cache inconsistencies

Immediate Consistency within a Cluster

- ✓ No stale reads
- ✓ No quorum reads

Workload Agnostic

- ✓ Read heavy
- ✓ Write heavy
- ✓ Mixed

Geographic Replication (XDR)

- ✓ Efficient Pipelining
- ✓ Asynchronous transfer

Hybrid Memory Architecture



- Flash Optimized Storage Layer: Indexes in DRAM, Data on SSD
- Multi-threaded & Massively Parallel
- **Self-healing Clusters**

Self-healing Clusters

- High Uptime
- Self-healing
- Low Management

Strong Consistency with High Availability

Guaranteed Strong Consistency

- ✓ Even with network splits

Rack-Awareness used

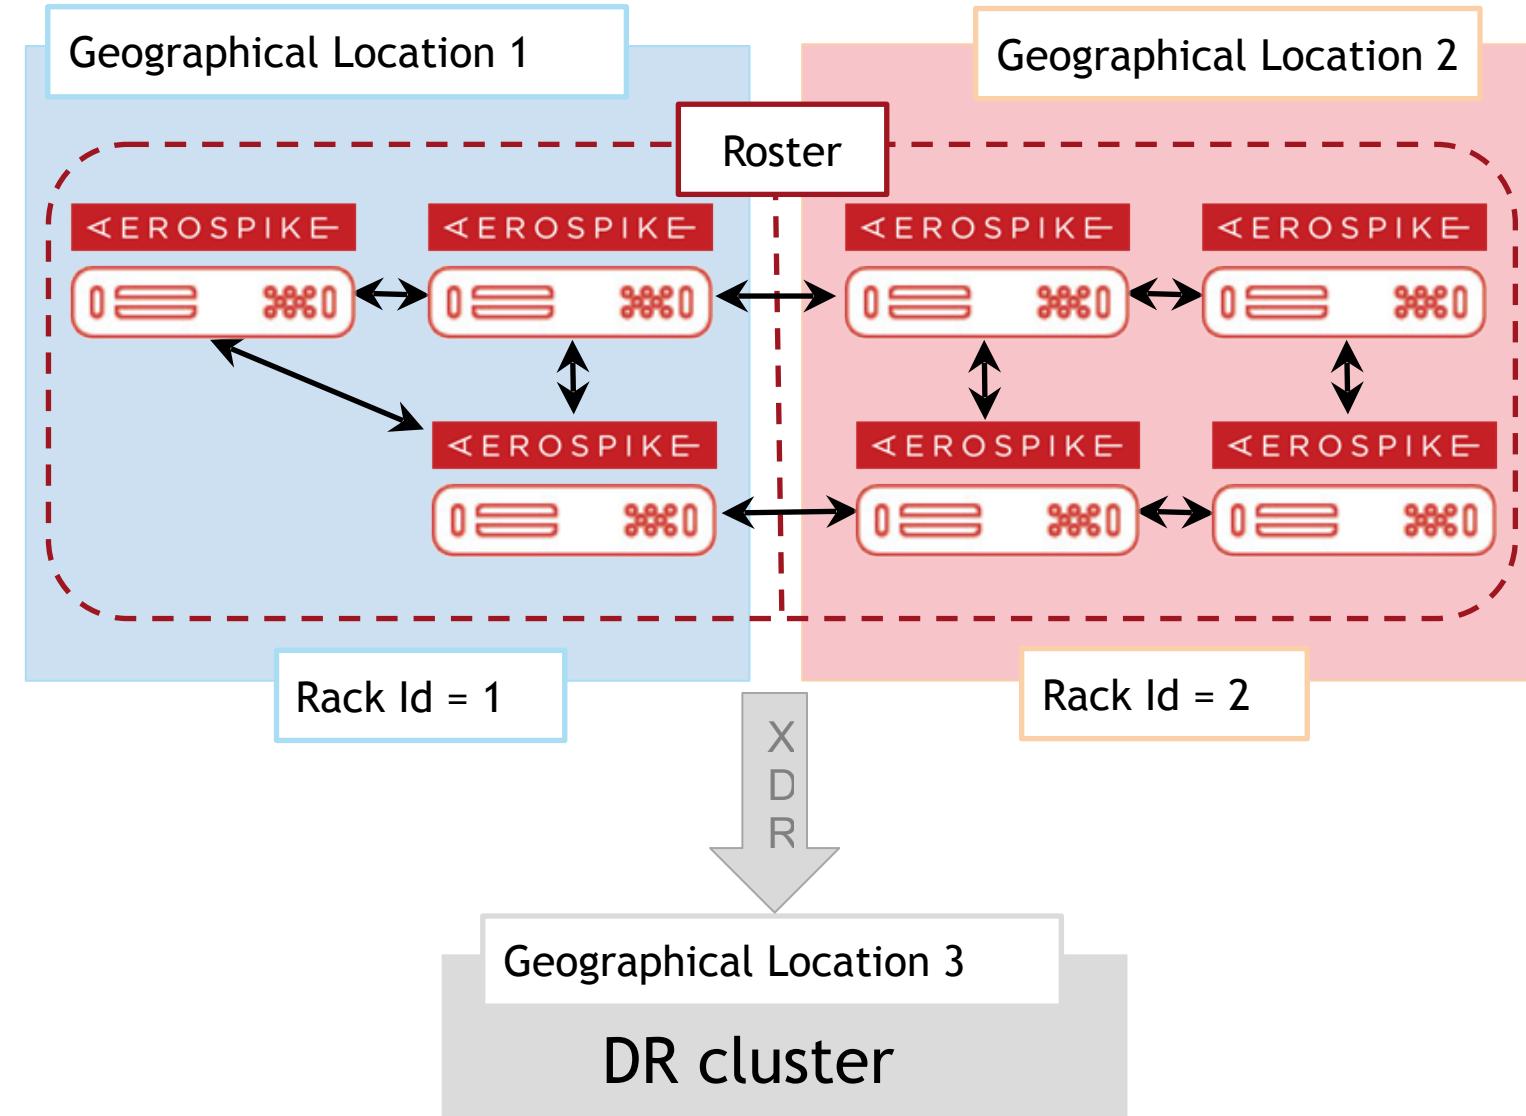
- ✓ Master and replica in different data centers
- ✓ Synchronous writes

High Availability

- ✓ Lose any node => 100% available
- ✓ Lose a DC => 100% available
- ✓ Might require one operator command

Low cost in Cloud

- ✓ Reads can use local DC
- ✓ Minor increase in chance of stale reads



Availability vs. Strong Consistency

“Aerospike does appear to provide linearizability through network partitions and process crashes”

- Kyle Kingsbury, Jepsen.io (<https://jepsen.io/analyses/aerospike-3-99-0-3>)

Availability vs. Strong Consistency

“Aerospike does appear to provide linearizability through network partitions and process crashes”

- Kyle Kingsbury, Jepsen.io (<https://jepsen.io/analyses/aerospike-3-99-0-3>)

	Linearizable Consistency	Sequential Consistency	Availability Mode
OPS (million)	1.87m	5.95m	6m
Read Latency	548 µs	225 µs	220 µs
Update Latency	630 µs	640 µs	640 µs

Configuration:

- In-memory
- Persistence enabled
- 5 node cluster
- 500M keys
- Replication factor 2
- 8 byte integer objects



Integrating Aerospike

Aerospike Clients

Find the latest Aerospike Client library for your favorite language.

Client libraries includes:

- Source or binary for the language
- Examples
- Benchmark tool

If you don't already have Aerospike Server up and running, [download and install it](#).

For more information on Aerospike's functionality, visit [the feature guides](#).

C / C++

4.6.13 C Client Library

[Download](#)

Java

4.4.9 Java Client Library

[Download](#)

Go

2.8.1 Go Client Library

[Download](#)

C# / .NET

3.9.3 C# Client Library

[Download](#)

Python

3.10.0 Python Client Library

[Download](#)

REST

1.1.1 Aerospike REST Client

[Download](#)

Node.js

3.14.1 Node.js Client Library

[Download](#)

Ruby

2.11.0 Ruby Client Library

[Download](#)

PHP

7.4.2 PHP Client Library

[Download](#)

Rust

0.4.0 Rust Client Library

[Download](#)

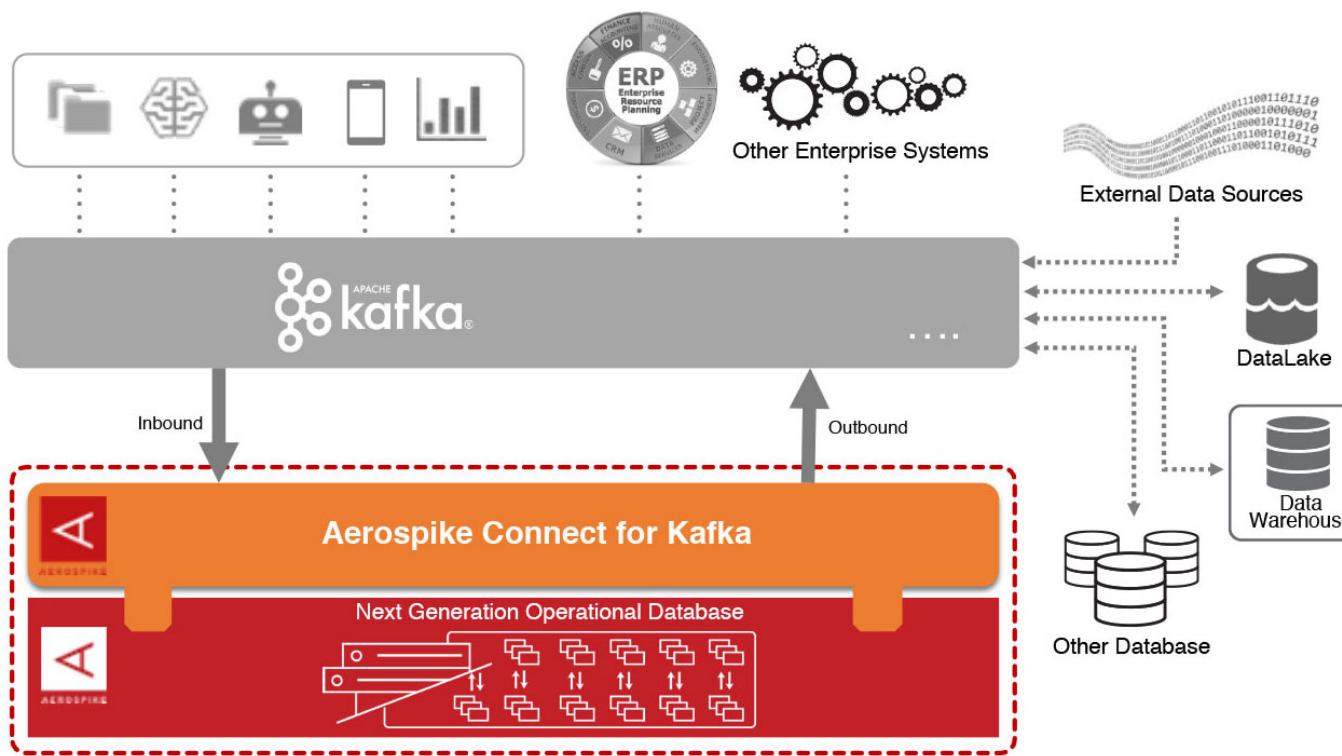
Orchestration Platforms



Pivotal **Cloud Foundry**



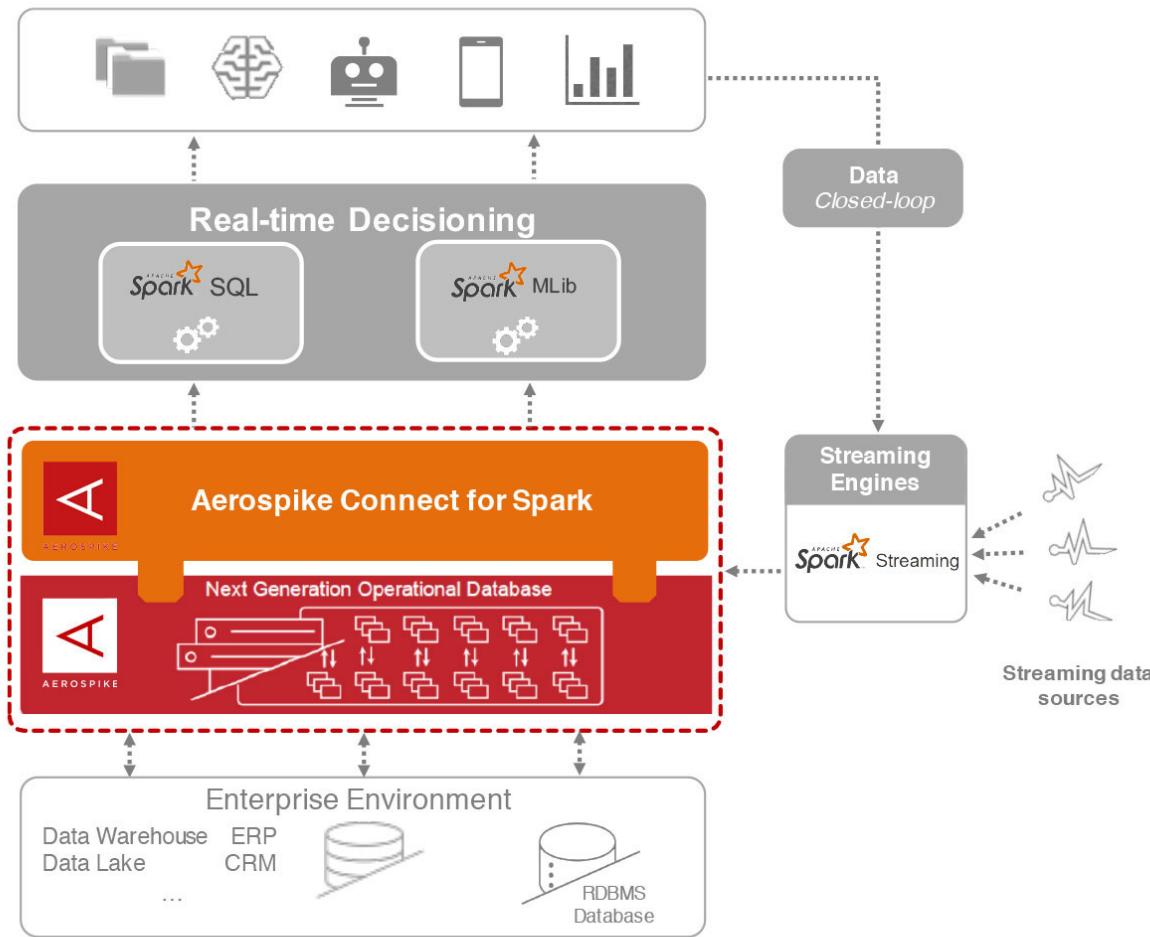
Connector for Kafka



Example Use Cases

- ✓ Edge-to-core data movement in fraud prevention systems and financial trading systems
- ✓ Edge-to-core data integration in retail user behavior at the POS, checkout with Product data for propensity & dynamic pricing
- ✓ Integration of operational data stored in Aerospike with DataLake / EDW

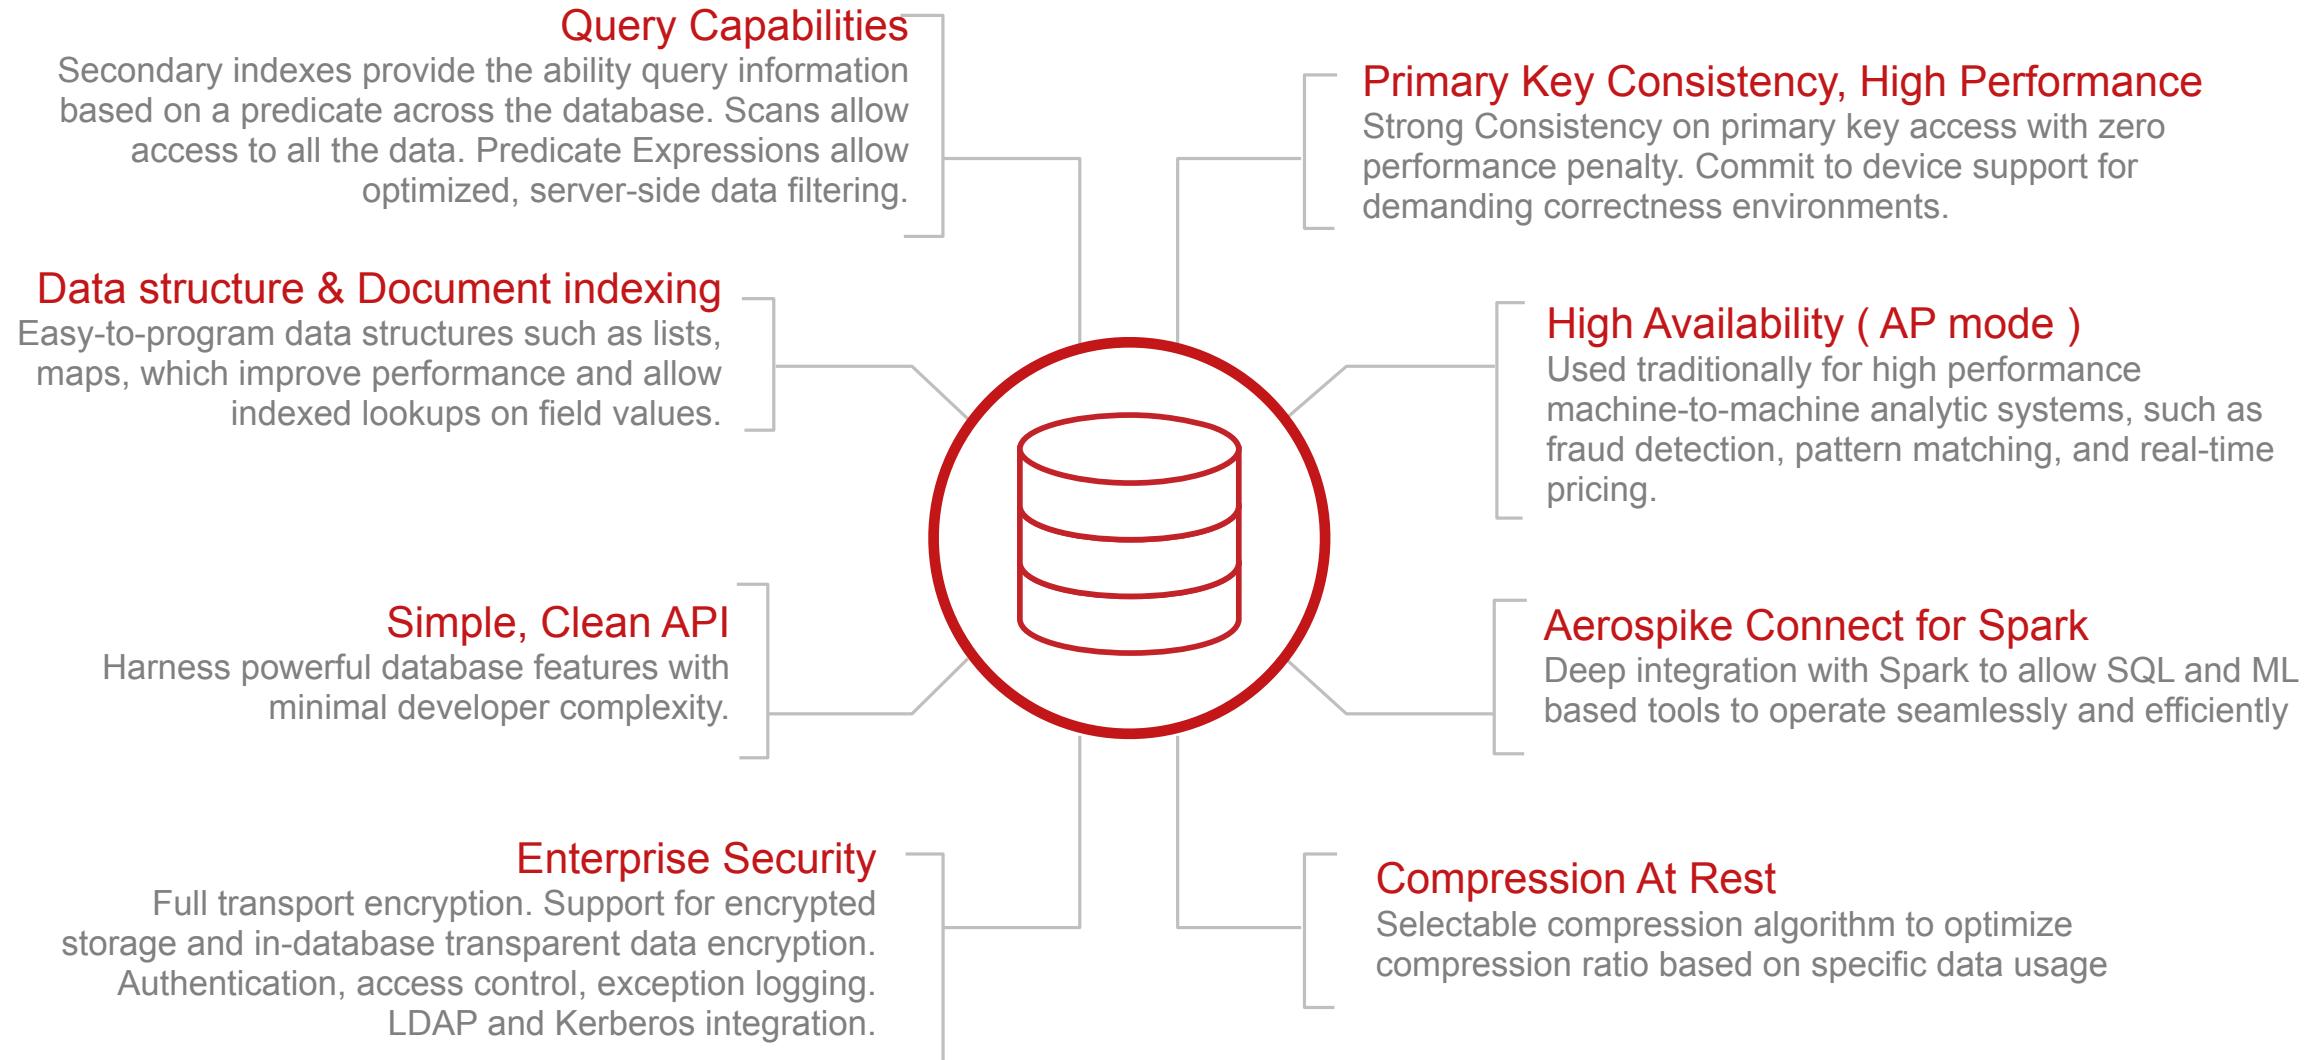
Connector for Spark



Example Use Cases

- ✓ **Fraud prevention:** transaction data via streaming and need to analyze based on historical data in real time
- ✓ **Recommendation Engines:** Real-time recommendations and targeting based on user behavior
- ✓ **Ad Tech:** Ad Fraud and real-time retargeting base on user behavior
- ✓ **Digital Identity Management**
- ✓ **Industrial Internet of Things (IIoT):** Real-time & closed loop business decisions

Other Features



Thank You!

Natalie Pistunovich
@NataliePis @aerospikeDB

